



# ***Network Service Delivery and Throughput Optimization Via Software Defined Networking***

*Thesis Defense:*

**Aaron Rosen**

Committee:

Kuang-Ching "KC" Wang  
Harlan Russell  
Sebastien Goasguen

Holcombe Department of Electrical &  
Computer Engineering

# Overview

---

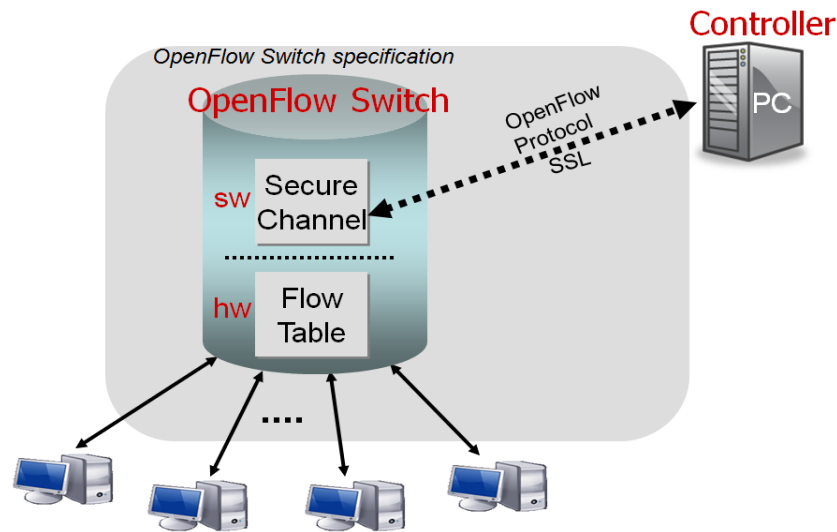
- **Network Service Delivery and Throughput Optimization Via Software Defined Networking**
  - Software Defined Networking (SDN)
  - Transmission Control Protocol (TCP)
  - Steroid OpenFlow Service (SOS)
  - GENI
  - Local Test bed

# Software Defined Networking

---

- Network architecture which decouples control plane and forwarding plane of network devices
- Allows network devices to be programmed via software API
- This allows new ideas to be quickly prototyped and tested at large scale
- OpenFlow is one example of SDN

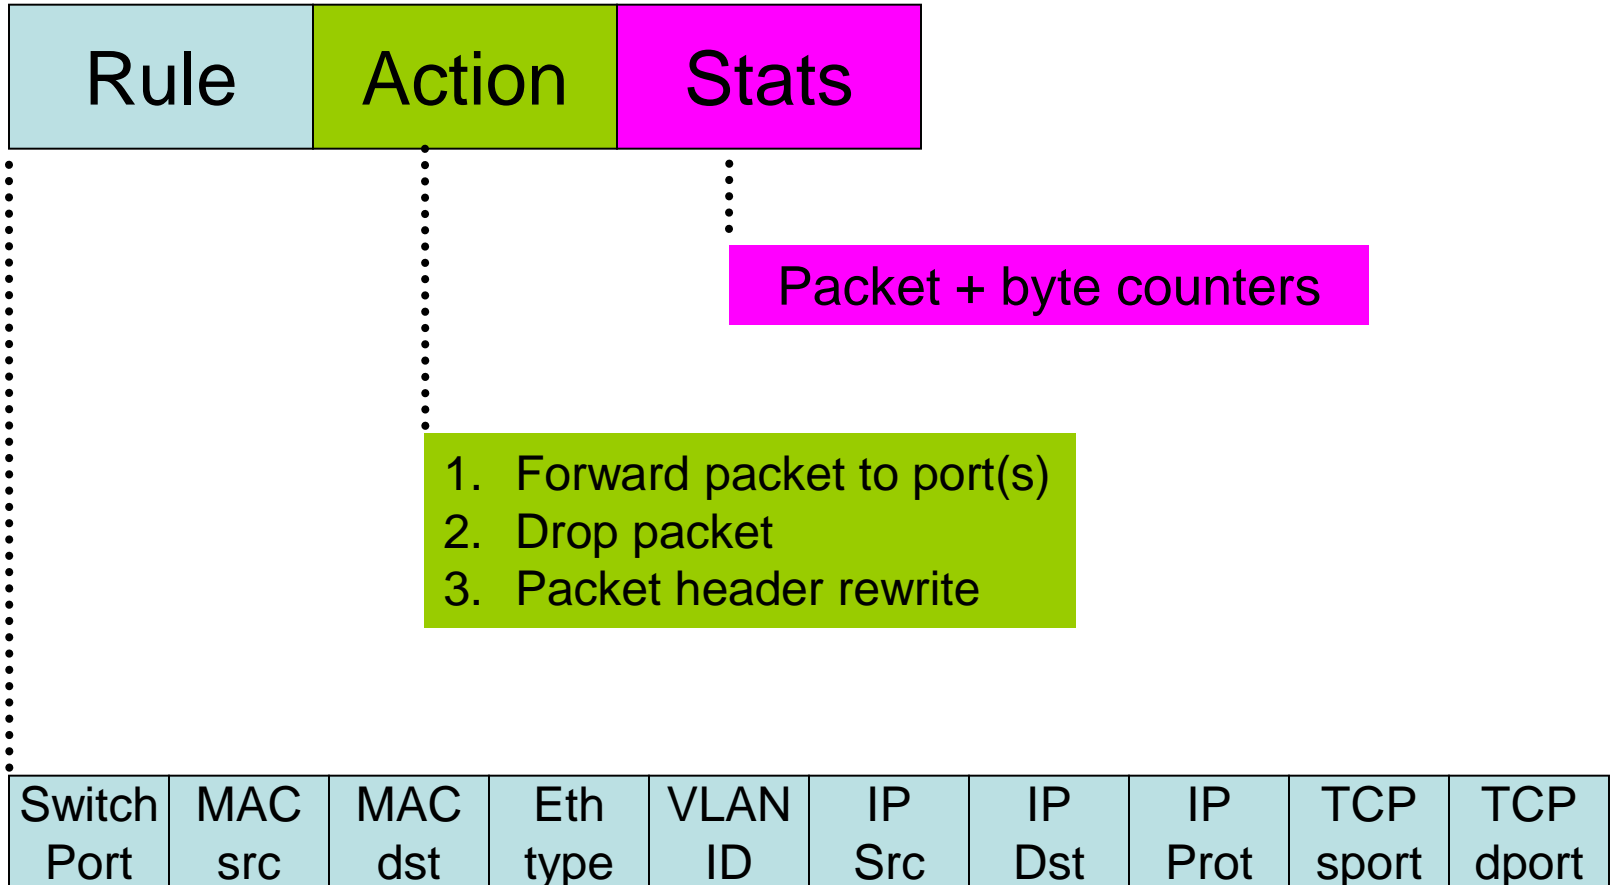
# The OpenFlow SDN Approach



*Illustration Courtesy of Stanford University  
Clean-Slate Program*


- Switches connect to remote controllers
- Controllers install flow entries into switches to handle packets
- Packets not matching any flow entry are sent to the controller (via Packet\_in)
- Controller decides how to handle packet
- Allows traffic to be manipulated easily

# OpenFlow Tables




# Transmission Control Protocol (TCP)

- De facto protocol for transmitting data reliably over the internet.
- Uses congestion control algorithms
- Performance degrades heavily with packet loss and high latency.
- High latency results in a higher Bandwidth Delay Product requiring a large buffer to store packets in flight before they are acknowledged




# Transmission Control Protocol (TCP)

- Lots of past work has investigated and proposed solutions to these issues.
- Unfortunately, most of these solutions require modifications to end users machines.
- GridFTP is one example that is used to move terabytes of data across the globe each year generated by the Large Hadron Collider.
- Requiring specialized software adds complexity and cost for end users.



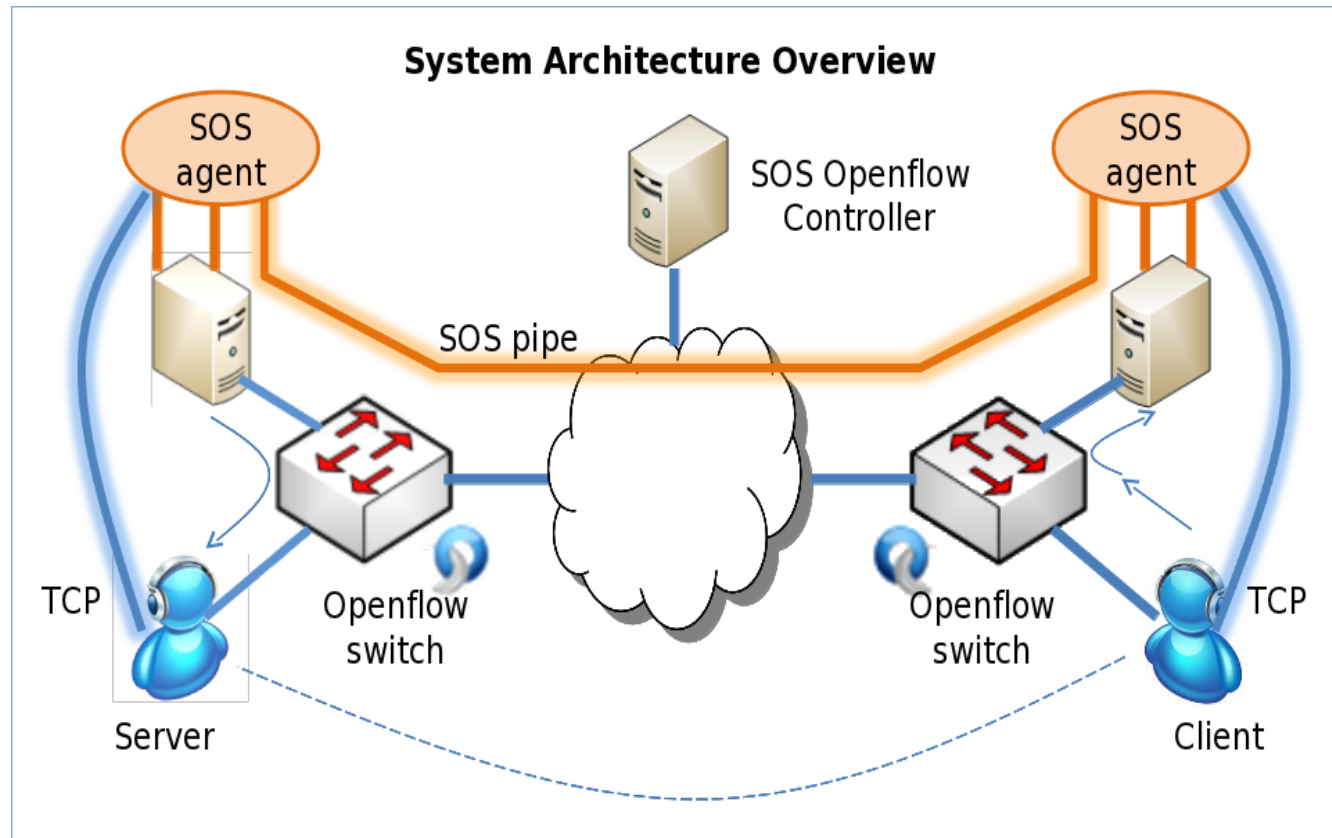
- Can these complexities be removed with SDN?





# Steroid OpenFlow Service

- SOS provides a seamless enhancement to end-to-end application throughput over long range networks.
- Decouples users protocol from network




# SOS Overview

- Goal: no modifications to host, seamless improvement
- Solution:
  - OpenFlow network detects TCP connection (client-server)
  - OpenFlow network redirects connection to local SOS agent
  - SOS agent starts high throughput transport to SOS agent on destination site
  - Destination SOS agent starts TCP connection to server
  - OpenFlow network discovers all sites with SOS agents
  - OpenFlow network allows multiple path transport

# Example Topology

- ..... OpenFlow Control Plane
- Physical Connection

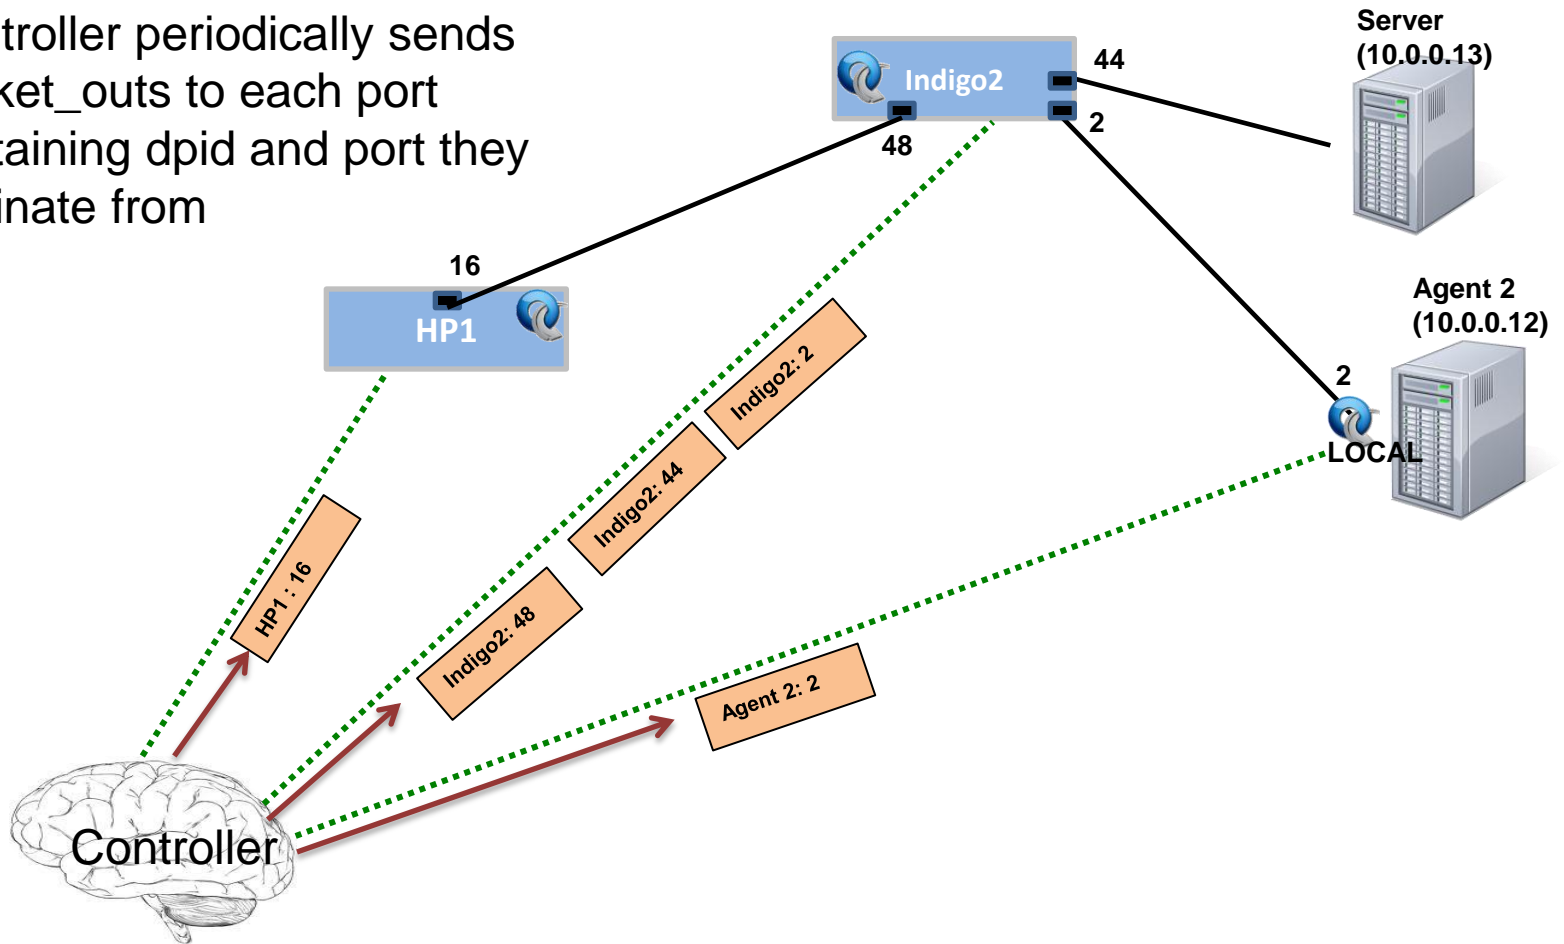


# Datapath Topology Discovery Controller Sends Packet\_out

..... OpenFlow Control Plane


— Physical Connection

Controller periodically sends packet\_outs to each port containing dpid and port they originate from



# Discovery packets leave desired interface

- ..... OpenFlow Control Plane
- Physical Connection

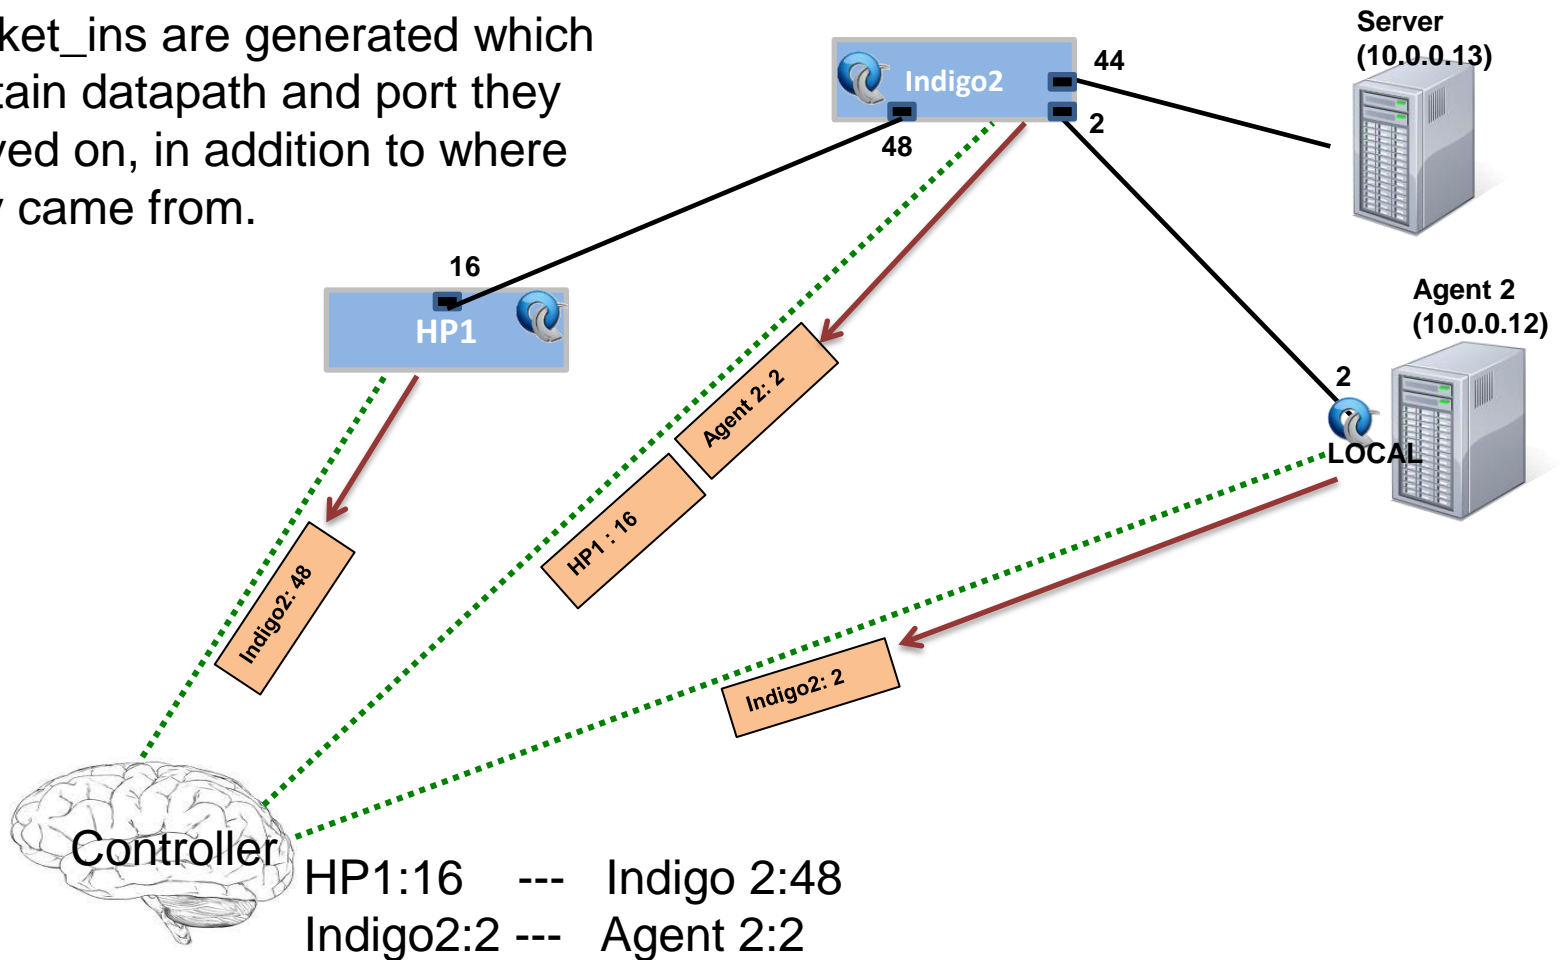


# Packets are returned to controller via Packet\_in

..... OpenFlow Control Plane

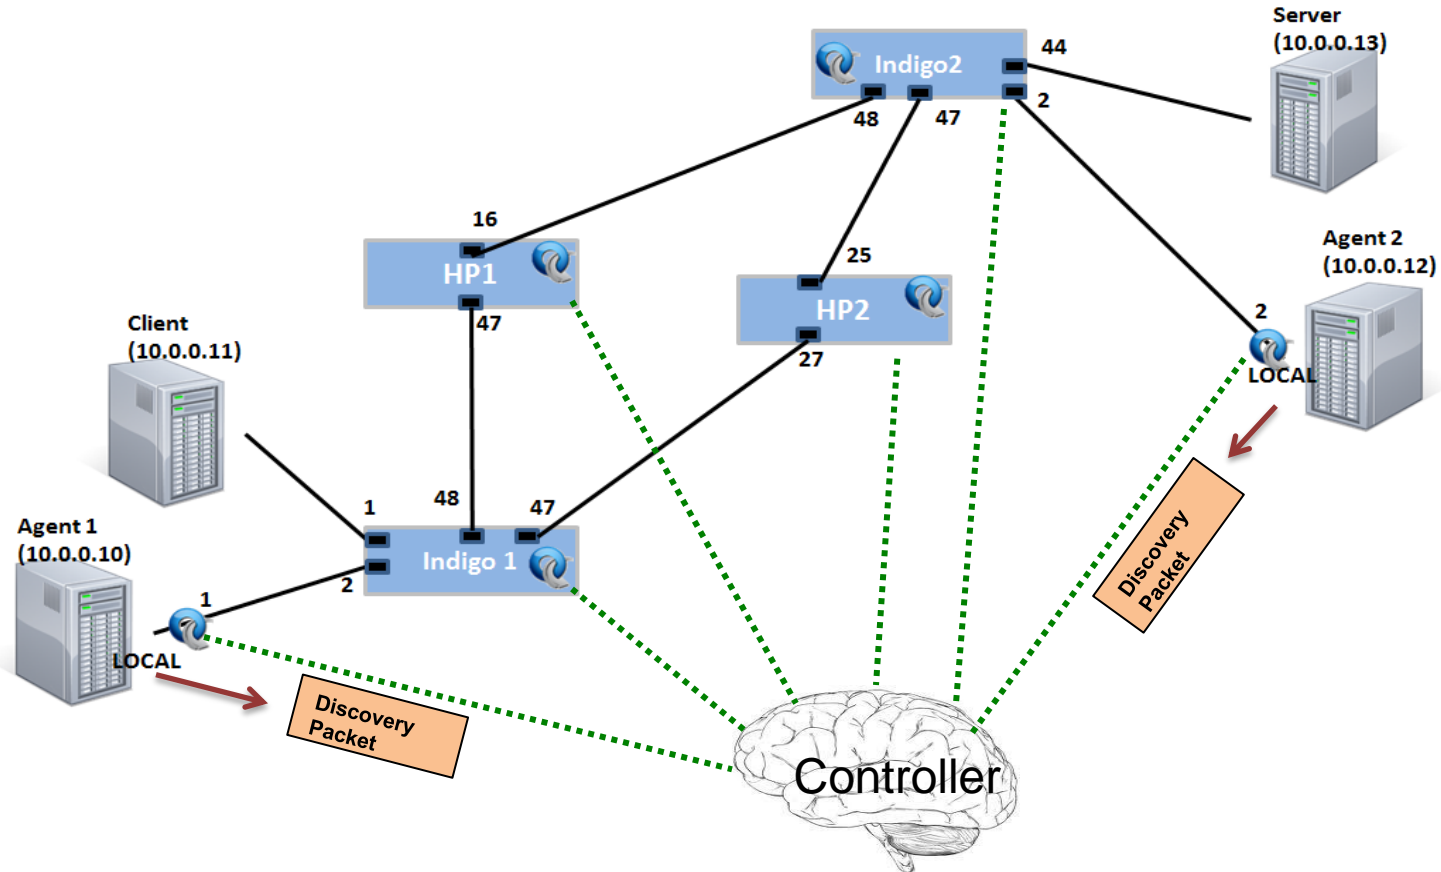
— Physical Connection

Packet\_ins are generated which contain datapath and port they arrived on, in addition to where they came from.

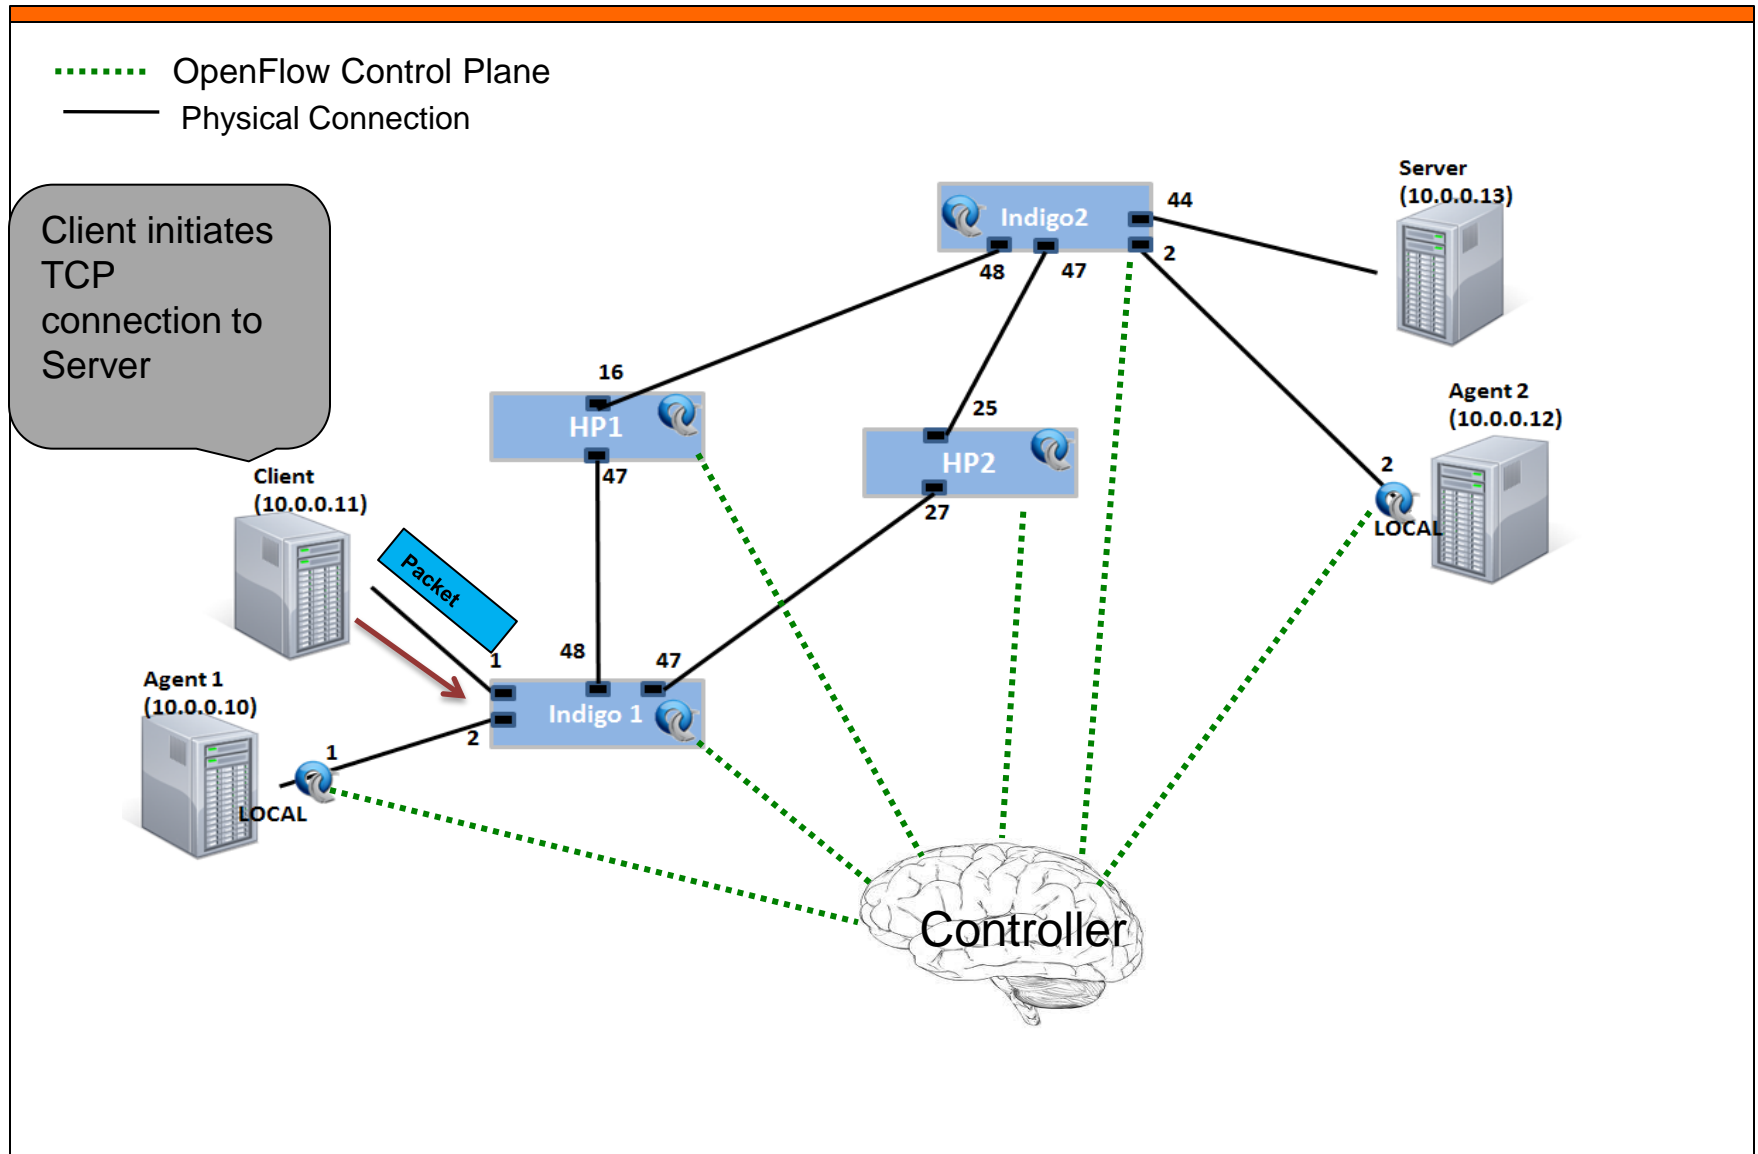


# Agent Discovery

- ..... OpenFlow Control Plane
- Physical Connection



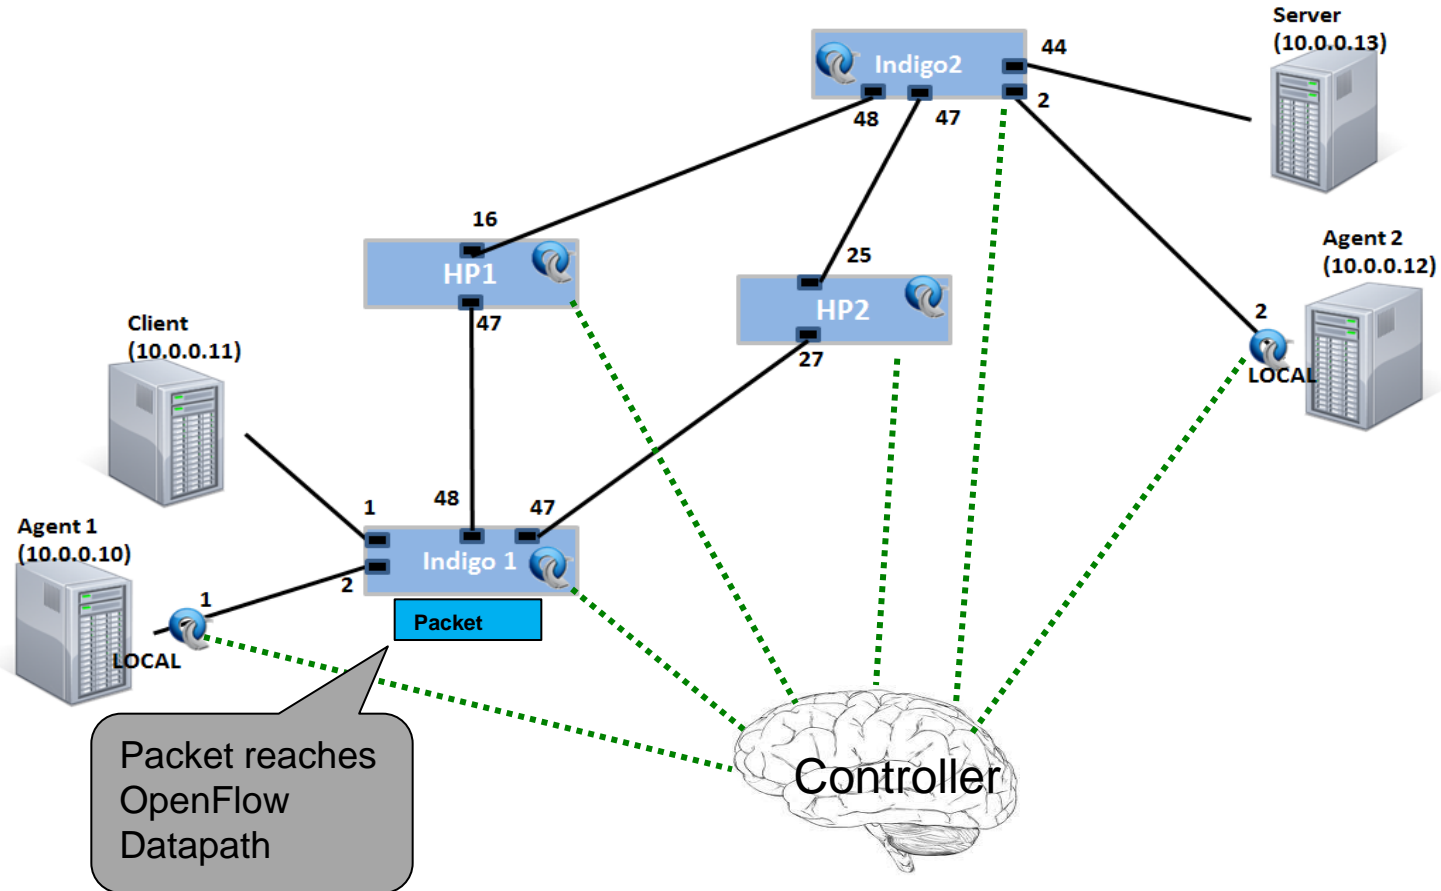
# SOS: Client Initiates TCP Connection






# Packet Reaches First OpenFlow Datapath

- ..... OpenFlow Control Plane
- Physical Connection



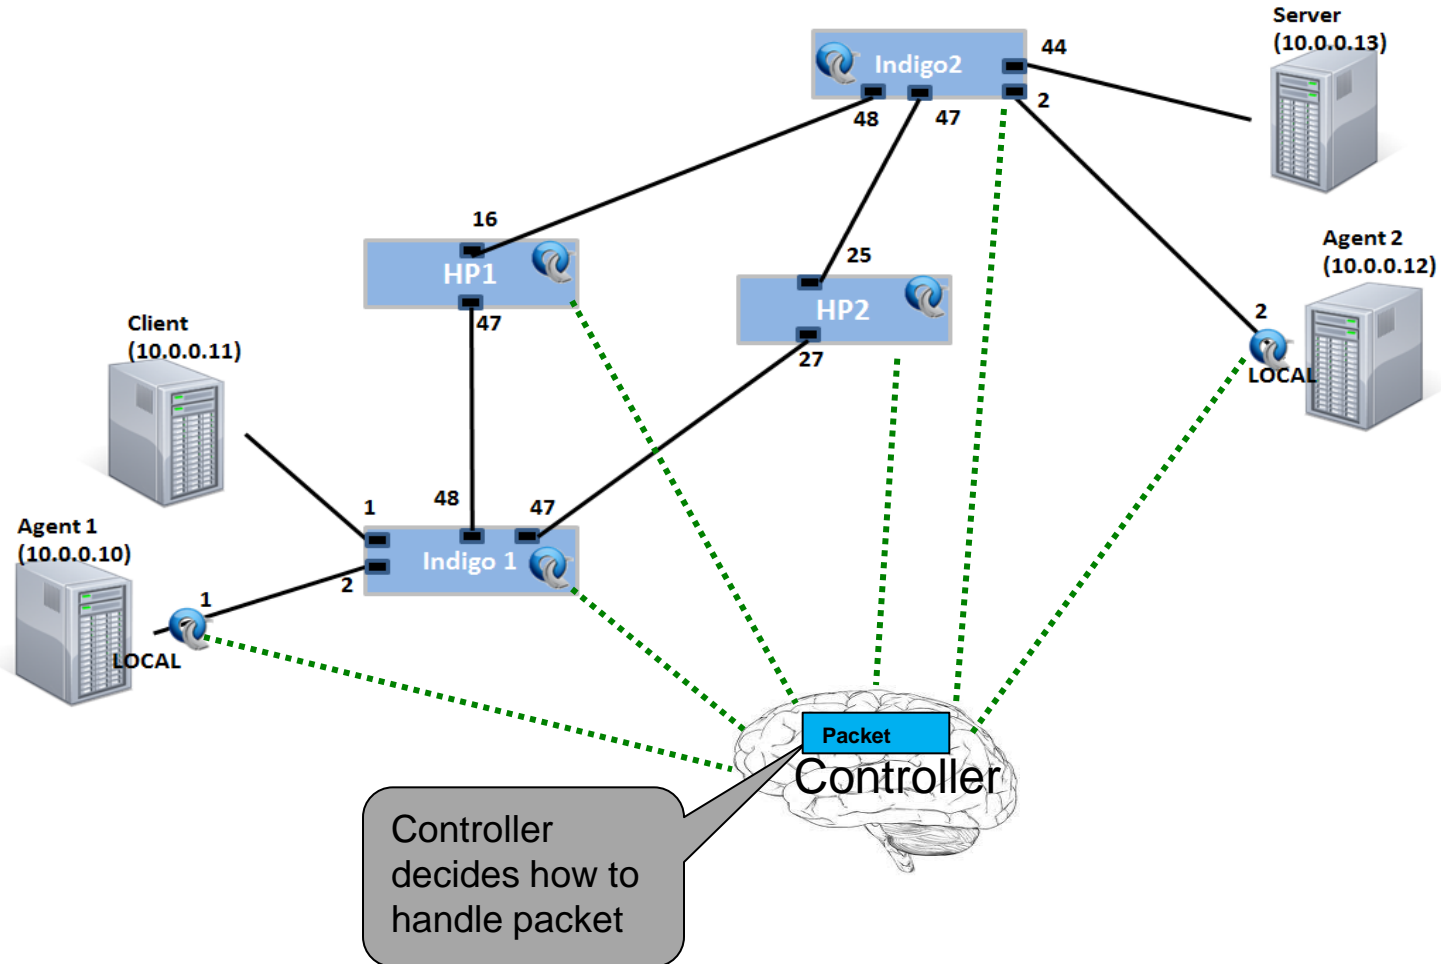
# Packet Forwarded to Controller

- ..... OpenFlow Control Plane
- Physical Connection



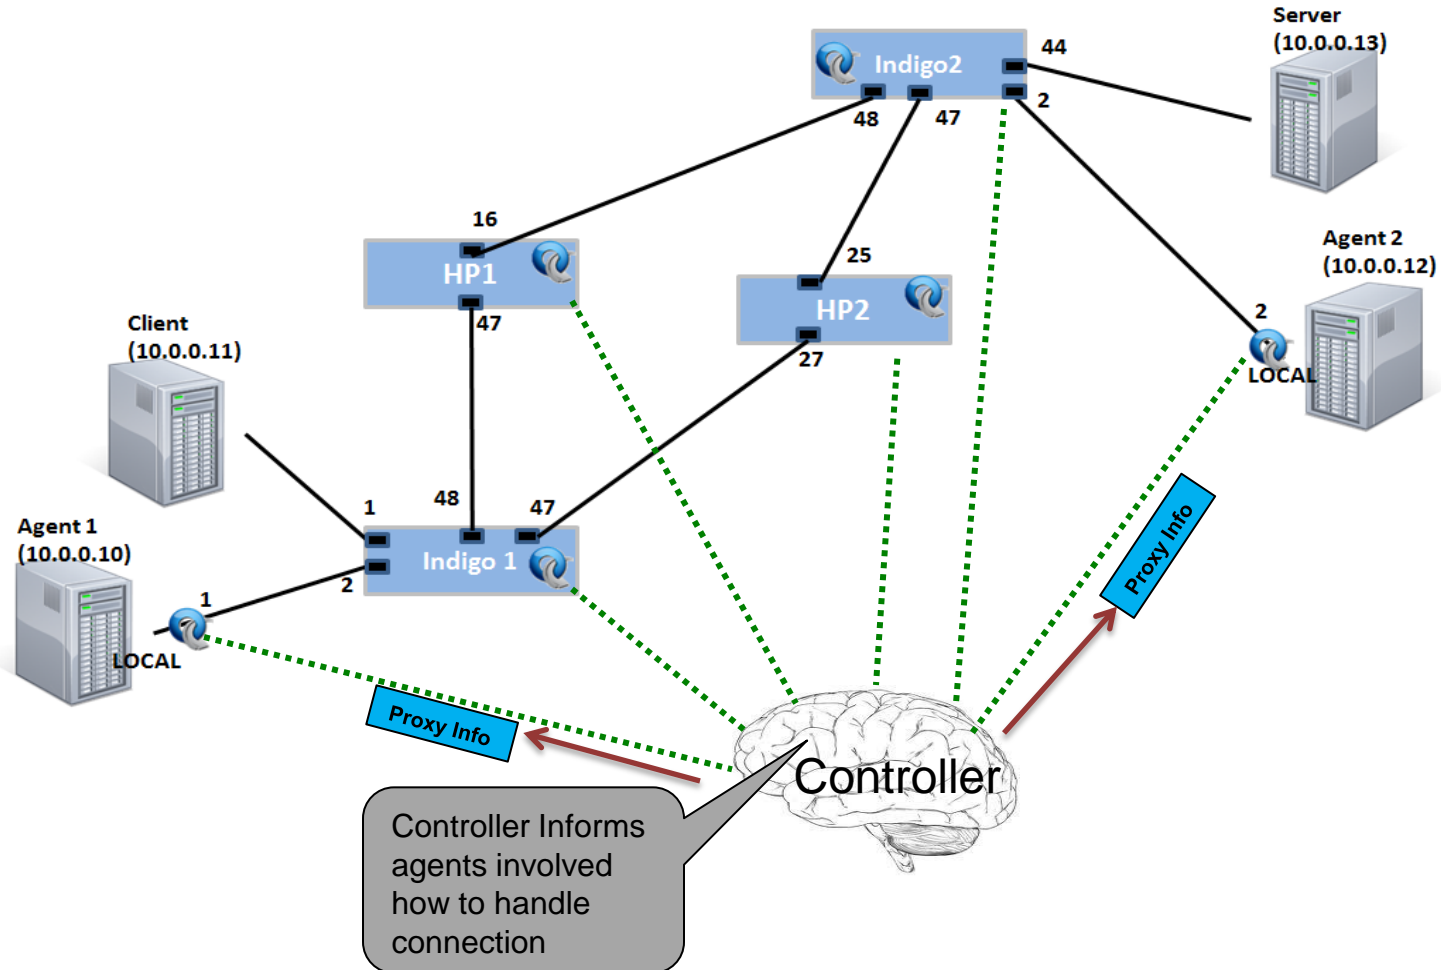
# Controller Handles Packet

- ..... OpenFlow Control Plane
- Physical Connection



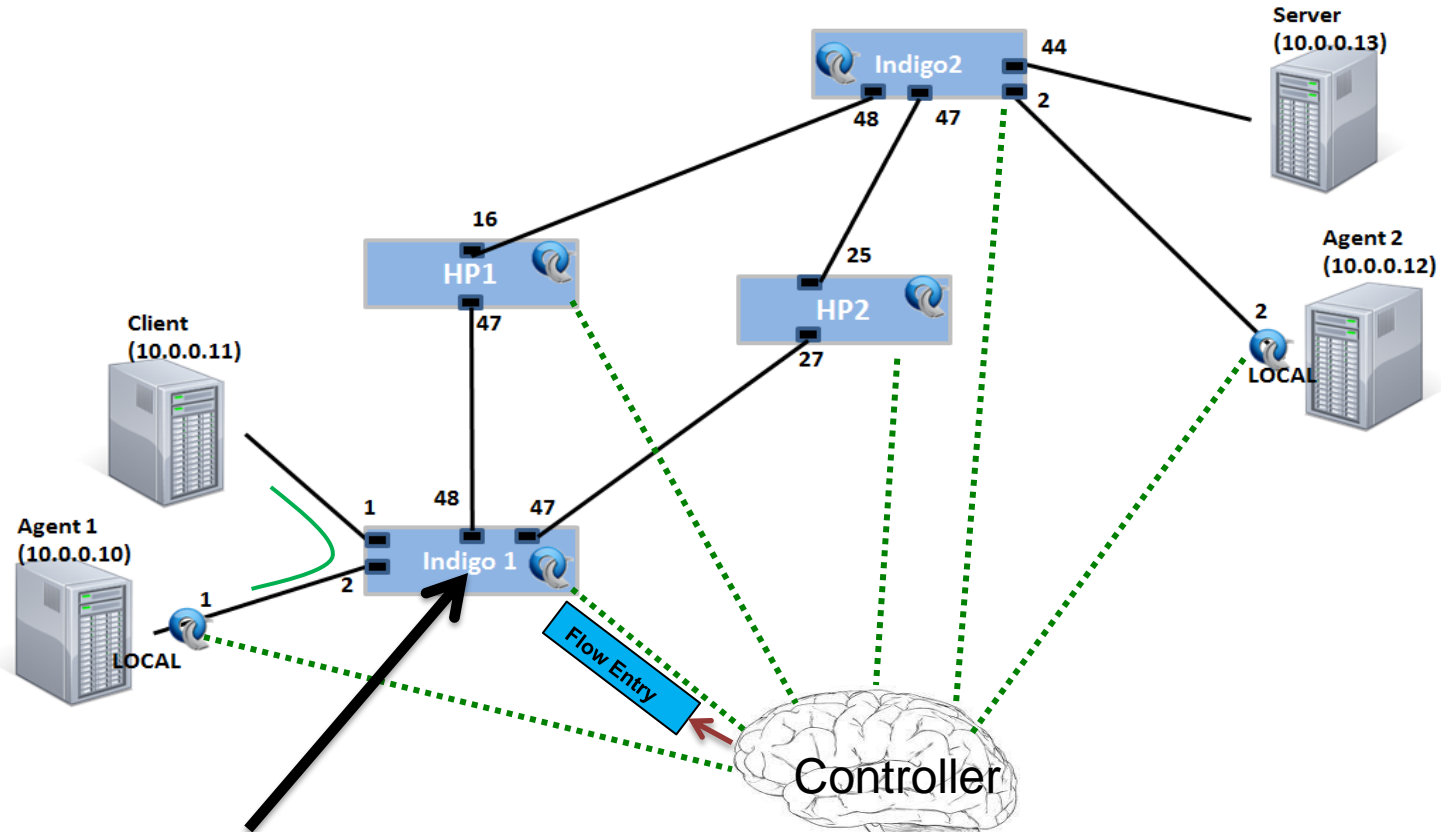
# Controller Informs Agents

- ..... OpenFlow Control Plane
- Physical Connection



# Redirecting traffic from Client to Agent 1

- ..... OpenFlow Control Plane
- Physical Connection

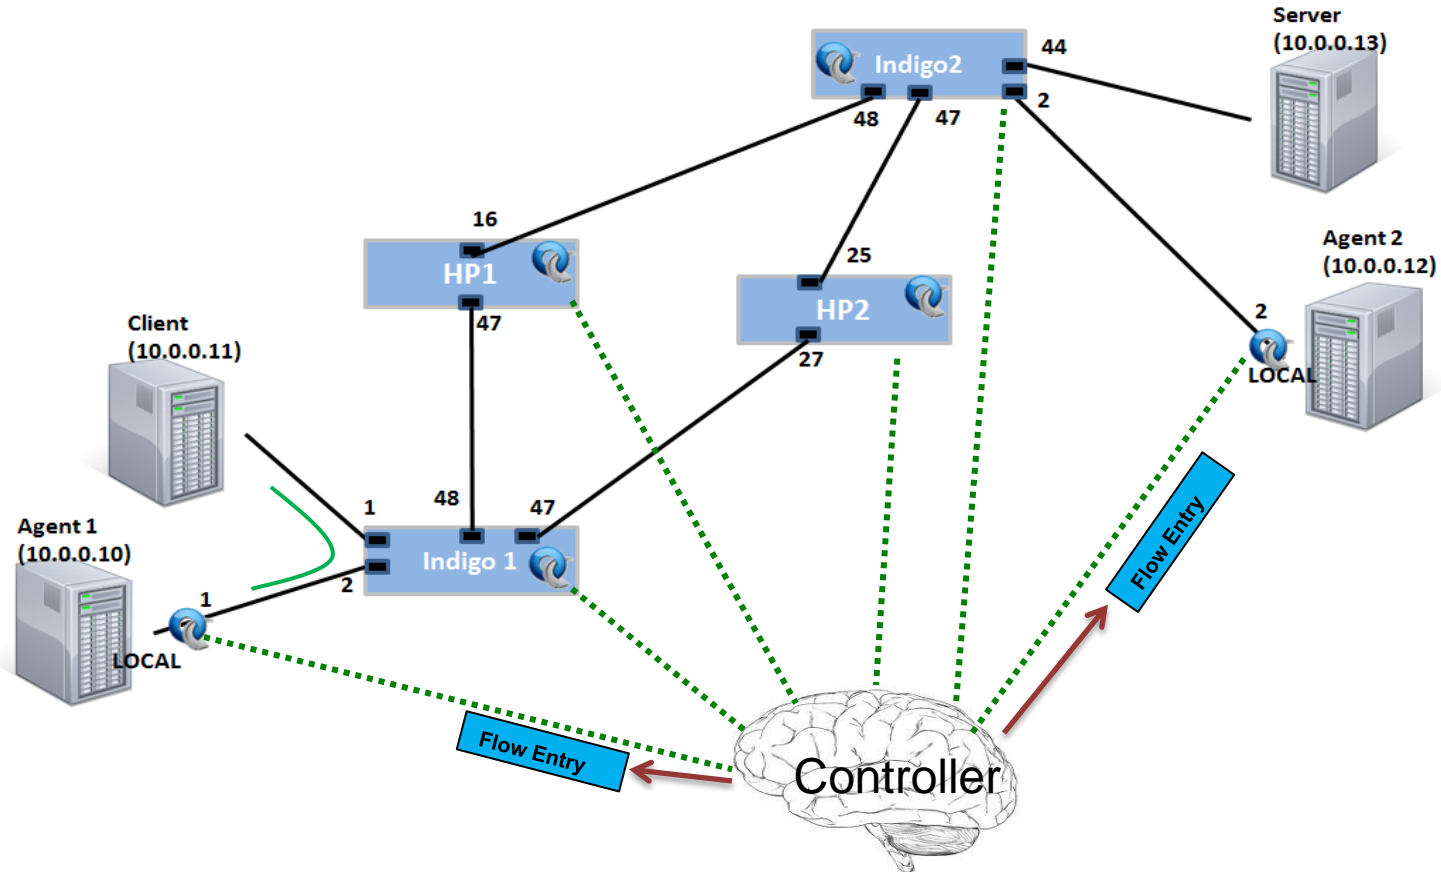



```
in_port=1,dl_vlan=0xffff,dl_vlan_pcp=0x00,dl_src=00:1f:29:32:91:99,dl_dst=00:1b:21:6a:85:88,nw_src=10.0.0.11,nw_dst=10.0.0.13,  
nw_tos=0x00,tp_src=41922,tp_dst=5003,actions=output:2
```

```
in_port=2,dl_dst=00:1f:29:32:91:99,nw_src=10.0.0.13,nw_dst=10.0.0.11,tp_src=5003,tp_dst=41922,actions=output:1
```

# More Flow Entries are installed

- ..... OpenFlow Control Plane
- Physical Connection

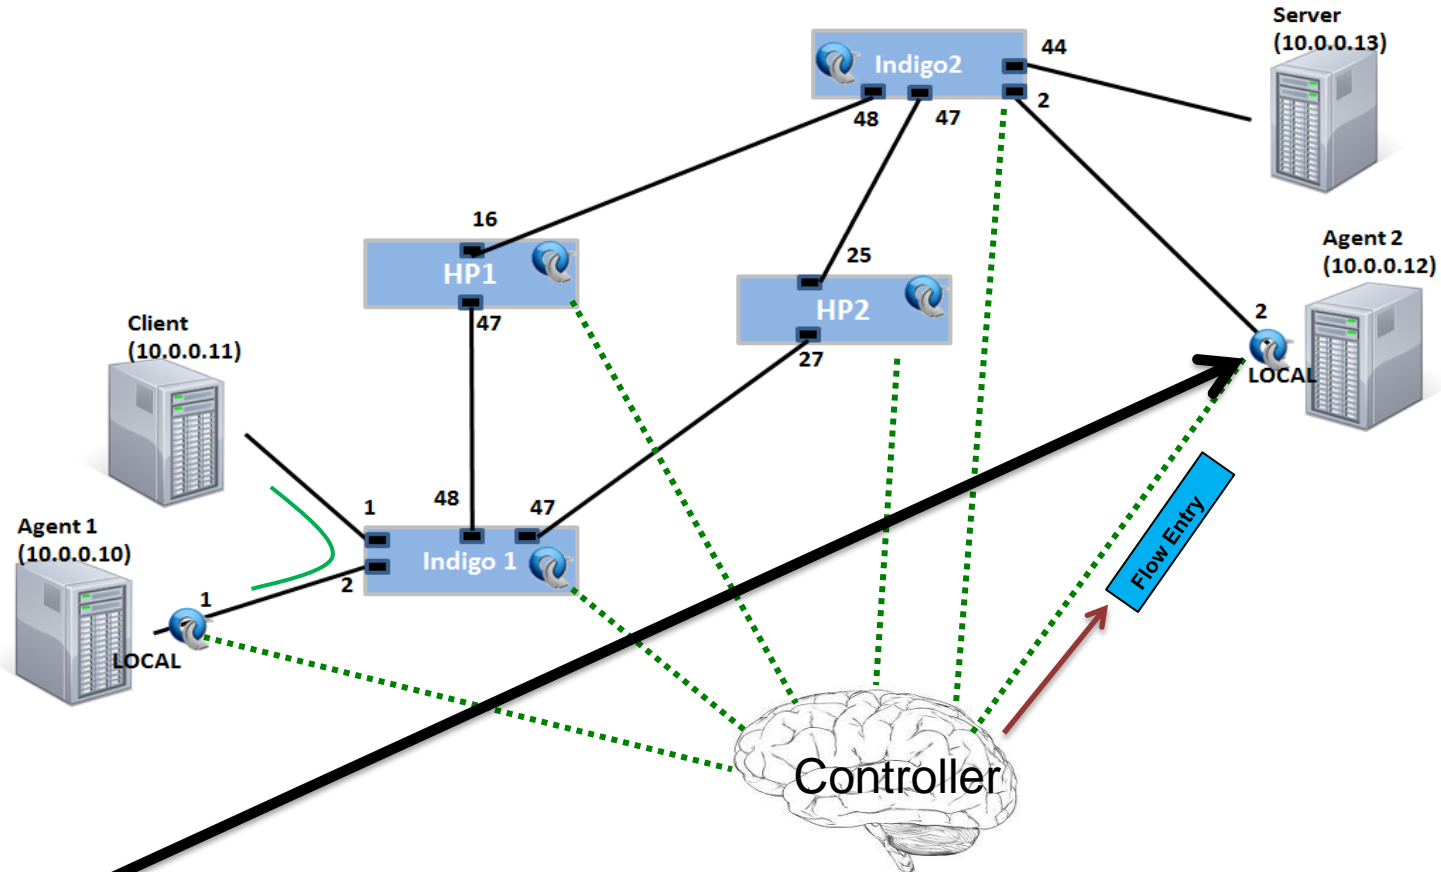




```
in_port=65534,dl_src=00:1f:29:32:92:4d,nw_src=10.0.0.10,nw_dst=10.0.0.11,tp_src=9877,tp_dst=41922
actions=mod_dl_src:00:1b:21:6a:85:88,mod_nw_src:10.0.0.13,mod_tp_src:5003,output:1
```

# Flow Entry Modifies Packet Headers for Server to Agent 2

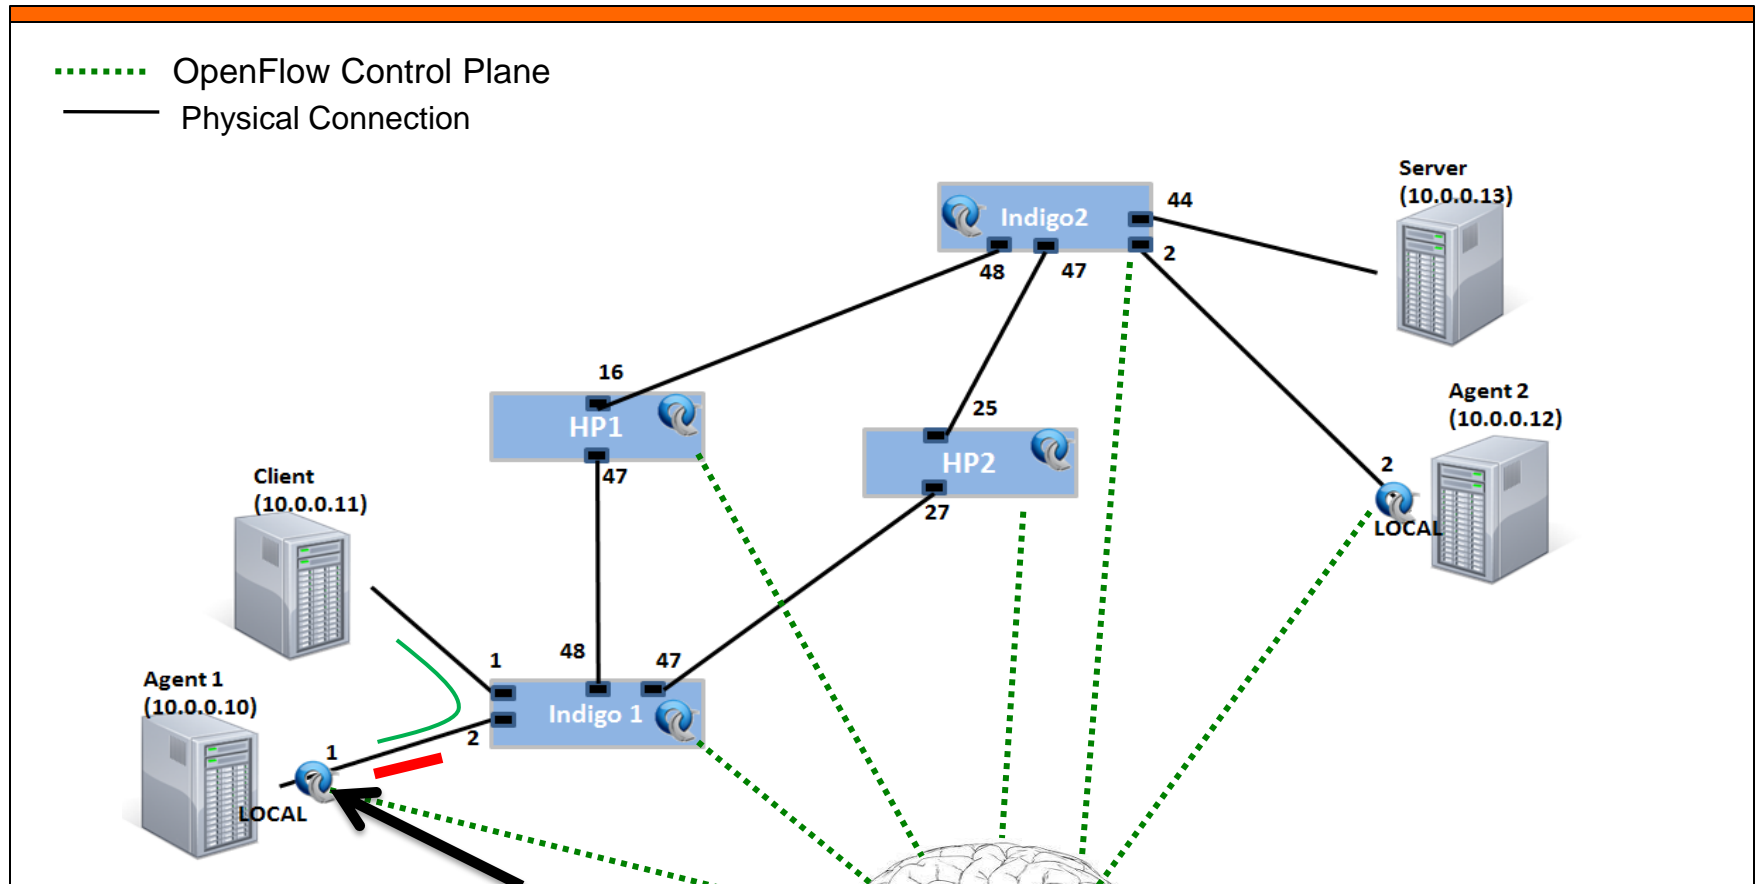
- ..... OpenFlow Control Plane
- Physical Connection



in\_port=2,nw\_src=10.0.0.13,nw\_dst=10.0.0.11,tp\_src=5003 actions=mod\_dl\_dst:00:1b:21:6b:50:df,mod\_nw\_dst:10.0.0.12,LOCAL



## Flow Entries on Agent 1 for Agent to Agent Communication



**Agent1 to Agent 2 flow entries using 4 streams and 2 different paths (see different mod\_dl\_src/dst)**

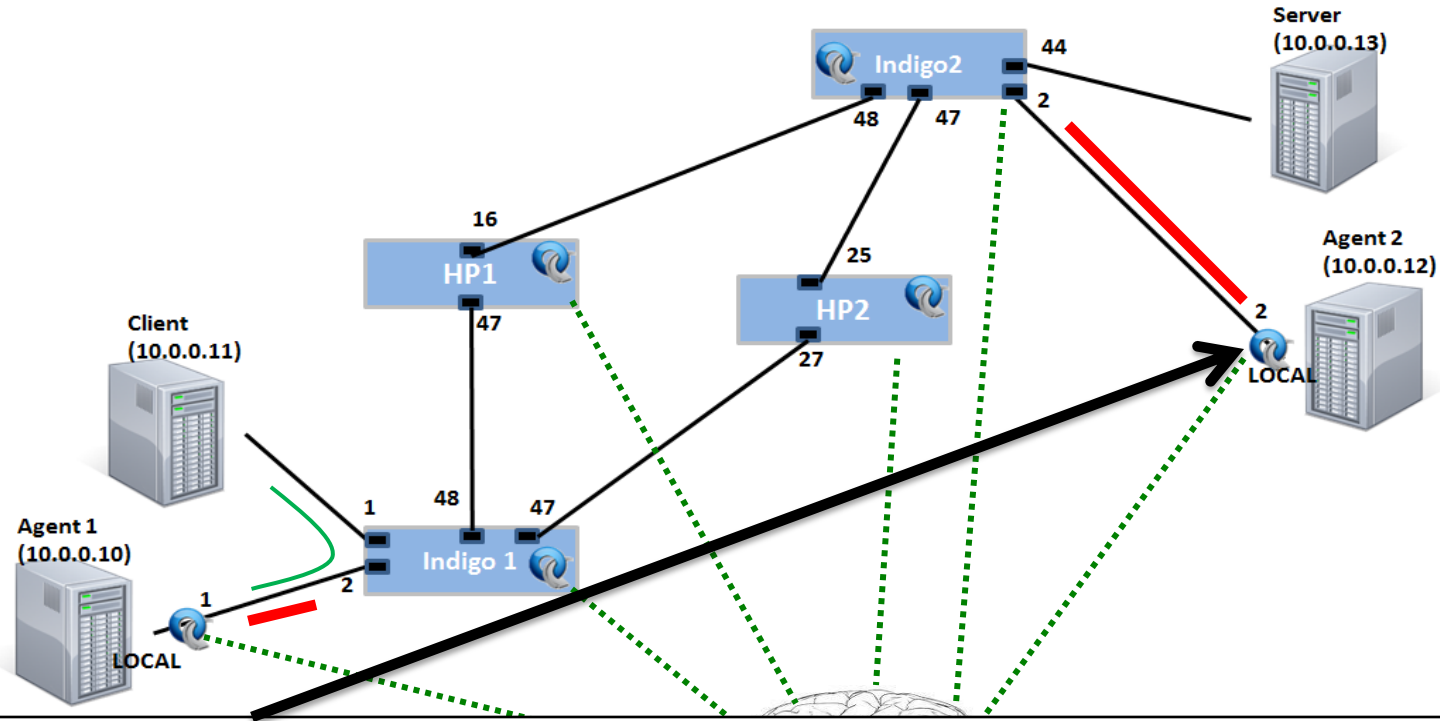
```
in_port=65534,nw_src=10.0.0.10,nw_dst=10.0.0.12,tp_dst=9881actions=mod_dl_src:00:1f:29:32:92:4f,mod_dl_dst:00:1b:21:6b:50:e1,output:1
in_port=65534,nw_src=10.0.0.10,nw_dst=10.0.0.12,tp_dst=9880actions=mod_dl_src:00:1f:29:32:92:4e,mod_dl_dst:00:1b:21:6b:50:e0,output:1
in_port=65534,nw_src=10.0.0.10,nw_dst=10.0.0.12,tp_dst=9879actions=mod_dl_src:00:1f:29:32:92:4f,mod_dl_dst:00:1b:21:6b:50:e1,output:1
in_port=65534,nw_src=10.0.0.10,nw_dst=10.0.0.12,tp_dst=9878actions=mod_dl_src:00:1f:29:32:92:4e,mod_dl_dst:00:1b:21:6b:50:e0,output:1
```

**Agent 2 to Agent 1 flow entry resetting the MAC addresses to be the correct values**

```
in port=1,nw src=10.0.0.12,nw dst=10.0.0.10 actions=mod dl src:00:1b:21:6b:50:df,mod dl dst:00:1f:29:32:92:4d,LOCAL
```

# Flow Entries on Agent 2 for Agent to Agent Communication

..... OpenFlow Control Plane  
 — Physical Connection



**Agent 2 to Agent 1 flow entries using 4 streams and 2 different paths (see different mod\_dl\_dst)**


```
in_port=65534,nw_src=10.0.0.12,nw_dst=10.0.0.10,tp_src=9880actions=mod_dl_src:00:1b:21:6b:50:e0,mod_dl_dst:00:1f:29:32:92:4e,output:2
in_port=65534,nw_src=10.0.0.12,nw_dst=10.0.0.10,tp_src=9878actions=mod_dl_src:00:1b:21:6b:50:e0,mod_dl_dst:00:1f:29:32:92:4e,output:2
in_port=65534,nw_src=10.0.0.12,nw_dst=10.0.0.10,tp_src=9881actions=mod_dl_src:00:1b:21:6b:50:e1,mod_dl_dst:00:1f:29:32:92:4f,output:2
in_port=65534,nw_src=10.0.0.12,nw_dst=10.0.0.10,tp_src=9879actions=mod_dl_src:00:1b:21:6b:50:e1,mod_dl_dst:00:1f:29:32:92:4f,output:2
```

**Agent 1 to Agent 2 flow entry resetting the MAC addresses to be the correct values**

```
in_port=2,nw_src=10.0.0.10,nw_dst=10.0.0.12 actions=mod_dl_src:00:1f:29:32:92:4d,mod_dl_dst:00:1b:21:6b:50:df, LOCAL
```

# Flow Entries on Core for Agent to Agent Communication


..... OpenFlow Control Plane  
 — Physical Connection



in_port=2,dl_src=00:1f:29:32:92:4f,dl_dst=00:1b:21:6b:50:e1,nw_src=10.0.0.10,nw_dst=10.0.0.12,actions=output:48
in_port=2,dl_src=00:1f:29:32:92:4e,dl_dst=00:1b:21:6b:50:e0,nw_src=10.0.0.10,nw_dst=10.0.0.12,actions=output:47
in_port=48,nw_src=10.0.0.12,nw_dst=10.0.0.10,actions=output:2
in_port=47,nw_src=10.0.0.12,nw_dst=10.0.0.10,actions=output:2

# Flow Entries on Core for Agent to Agent Communication

..... OpenFlow Control Plane  
— Physical Connection

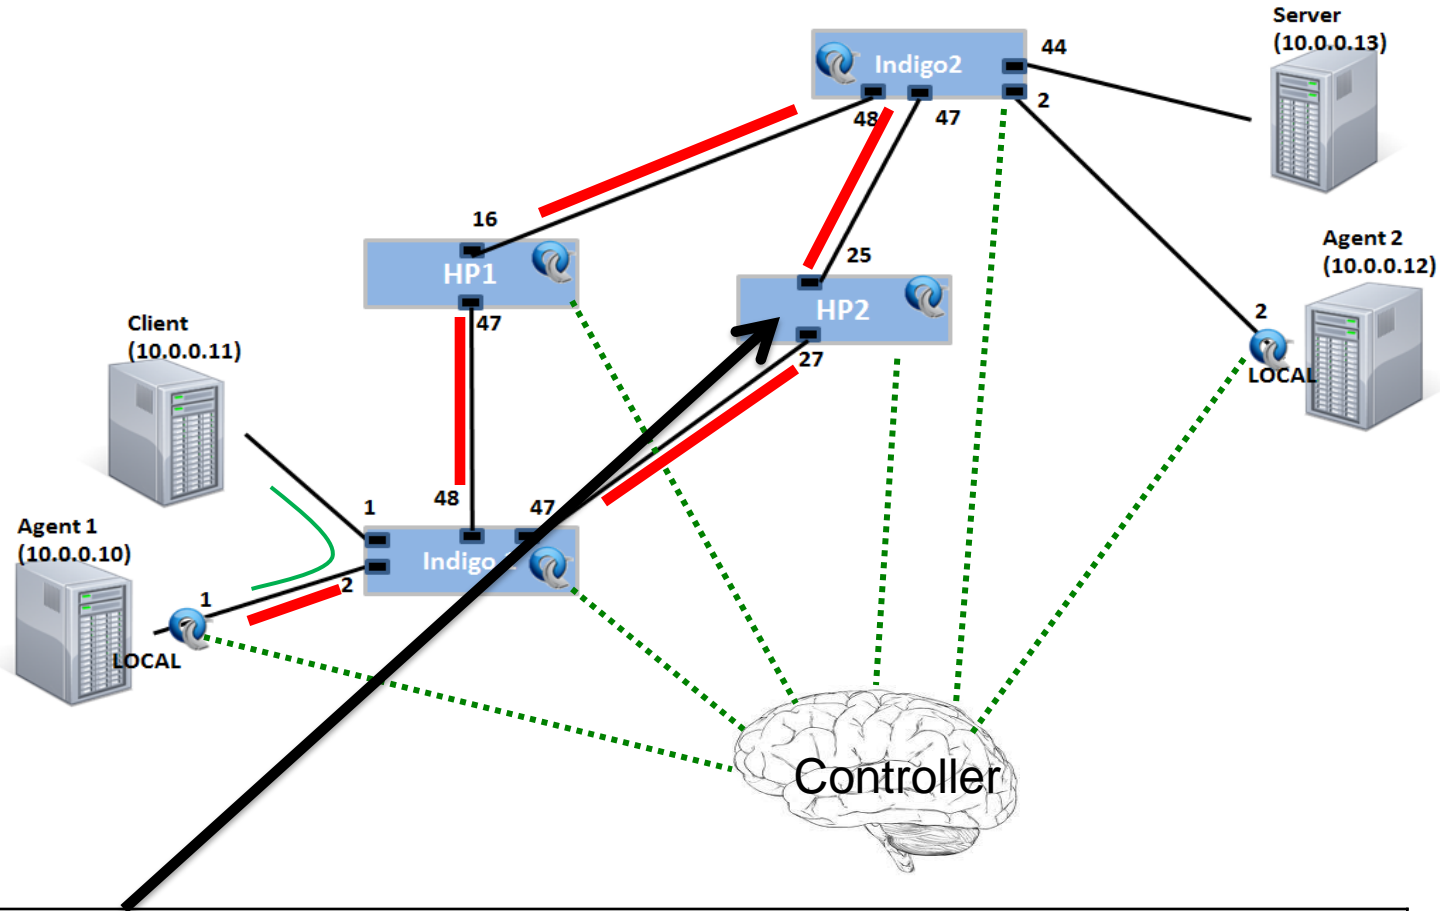


in\_port=16,dl\_src=00:1b:21:6b:50:e1,dl\_dst=00:1f:29:32:92:4f,nw\_src=10.0.0.12,nw\_dst=10.0.0.10,actions=output:47

in\_port=47,dl\_src=00:1f:29:32:92:4f,dl\_dst=00:1b:21:6b:50:e1,nw\_src=10.0.0.10,nw\_dst=10.0.0.12,actions=output:16

# Flow Entries on Core for Agent to Agent Communication

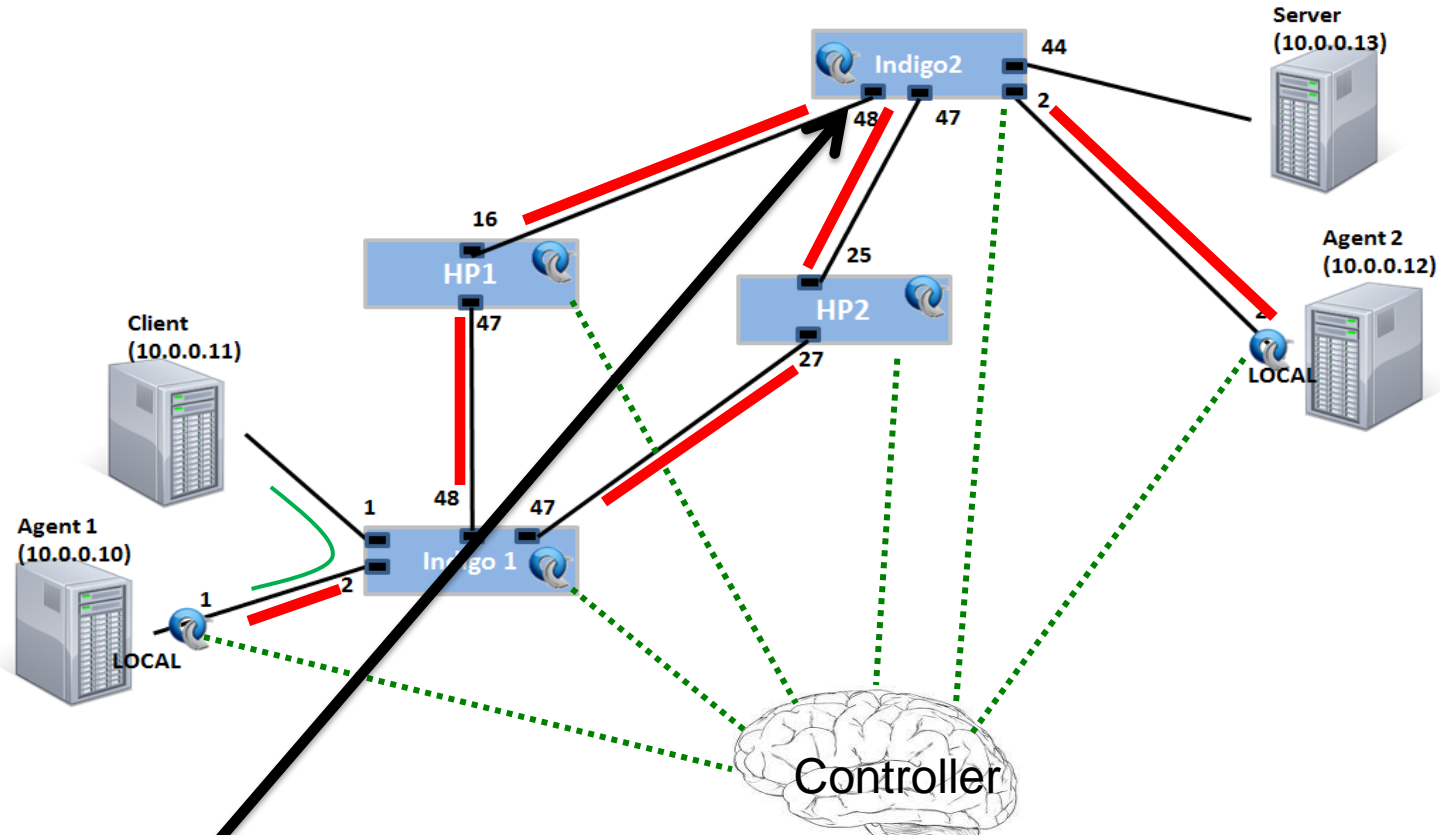
- ..... OpenFlow Control Plane
- Physical Connection



in_port=25,dl_src=00:1b:21:6b:50:e0,dl_dst=00:1f:29:32:92:4e,nw_src=10.0.0.12,nw_dst=10.0.0.10,actions=output:27
in_port=27,dl_src=00:1f:29:32:92:4e,dl_dst=00:1b:21:6b:50:e0,nw_src=10.0.0.10,nw_dst=10.0.0.12,actions=output:25

# Flow Entries on Core for Agent to Agent Communication


..... OpenFlow Control Plane  
——— Physical Connection



in_port=2,dl_src=00:1b:21:6b:50:e1,dl_dst=00:1f:29:32:92:4f,nw_src=10.0.0.12,nw_dst=10.0.0.10,actions=output:48
in_port=2,dl_src=00:1b:21:6b:50:e0,dl_dst=00:1f:29:32:92:4e,nw_src=10.0.0.12,nw_dst=10.0.0.10,actions=output:47
in_port=48,nw_src=10.0.0.10,nw_dst=10.0.0.12,actions=output:2
in_port=47,nw_src=10.0.0.10,nw_dst=10.0.0.12,actions=output:2

# Flow Entries to Forward Between Server and Agent 2

- ..... OpenFlow Control Plane
- Physical Connection

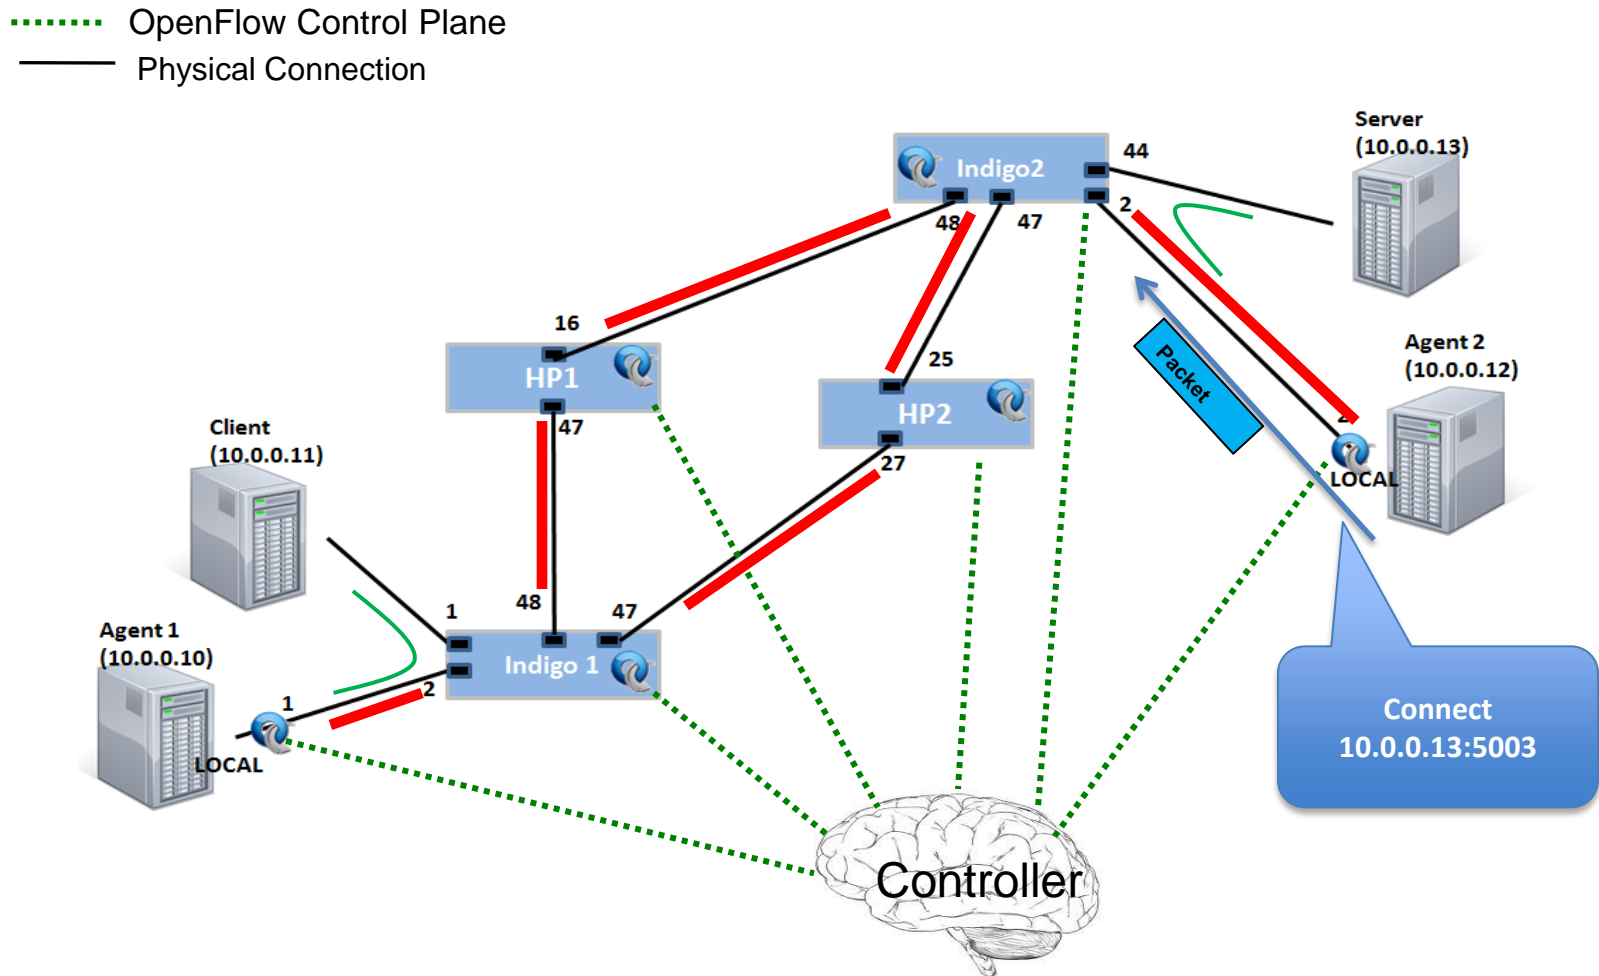


in\_port=2,nw\_src=10.0.0.11,nw\_dst=10.0.0.13,tp\_dst=5003,actions=output:44

in\_port=44,nw\_src=10.0.0.13,nw\_dst=10.0.0.11,tp\_src=5003,actions=output:2

The diagram illustrates a network topology with the following components and connections:

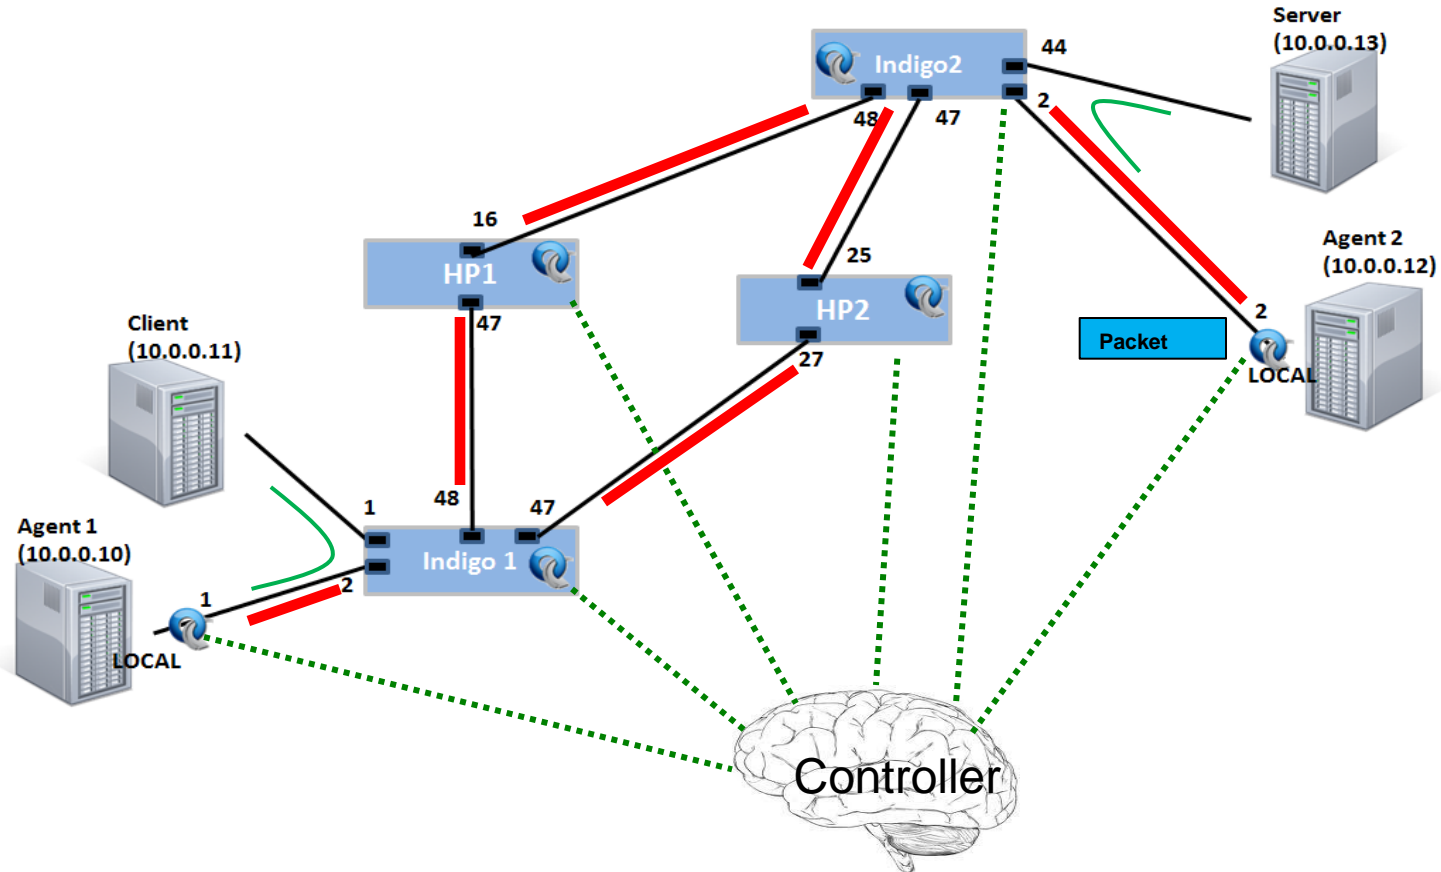
- Legend:**
  - OpenFlow Control Plane (Dotted green line)
  - Physical Connection (Solid black line)
- Devices and Connections:**
  - Indigo 1:** Connected to Indigo 2 (48 to 47), HP1 (16 to 47), and HP2 (25 to 27). It has a LOCAL interface connected to Agent 1 (10.0.0.10) and a physical connection to Indigo 2 (48 to 47).
  - Indigo 2:** Connected to Indigo 1 (47 to 48), HP1 (47 to 16), HP2 (47 to 25), and a Server (10.0.0.13) (44 to 2). It has a LOCAL interface connected to Agent 2 (10.0.0.12).
  - HP1:** Connected to Indigo 1 (47 to 16) and Indigo 2 (16 to 47).
  - HP2:** Connected to Indigo 1 (27 to 25) and Indigo 2 (25 to 47).
  - Agent 1 (10.0.0.10):** Connected to Indigo 1 (1 to 47) via a LOCAL interface.
  - Agent 2 (10.0.0.12):** Connected to Indigo 2 (2 to 44) via a LOCAL interface.
  - Server (10.0.0.13):** Connected to Indigo 2 (44 to 2).
  - Client (10.0.0.11):** Connected to Indigo 1 (1 to 47) via a LOCAL interface.
- Controller:** A central brain icon connected to all Indigo devices (Indigo 1 and Indigo 2) via the OpenFlow Control Plane (dotted green lines).
- Packet Flow:** A blue arrow labeled "Packet" points from the Server towards the Indigo 2 LOCAL interface.
- Callout:** A blue box contains the text "Connect 10.0.0.13:5003".





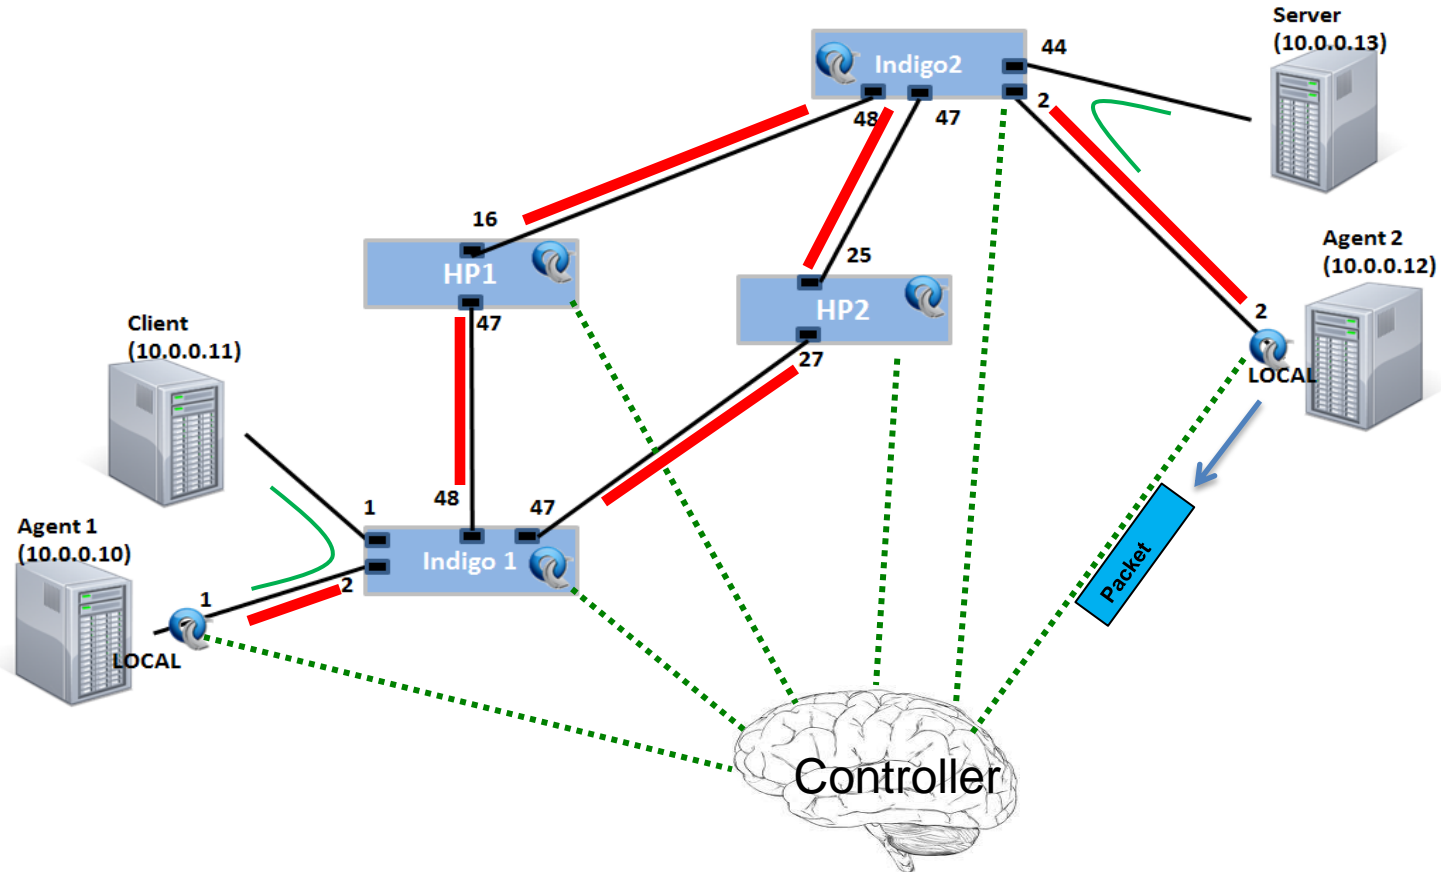
# Packet reaches OpenFlow Datapath

- ..... OpenFlow Control Plane
- Physical Connection




# Forwarded to controller

- ..... OpenFlow Control Plane
- Physical Connection



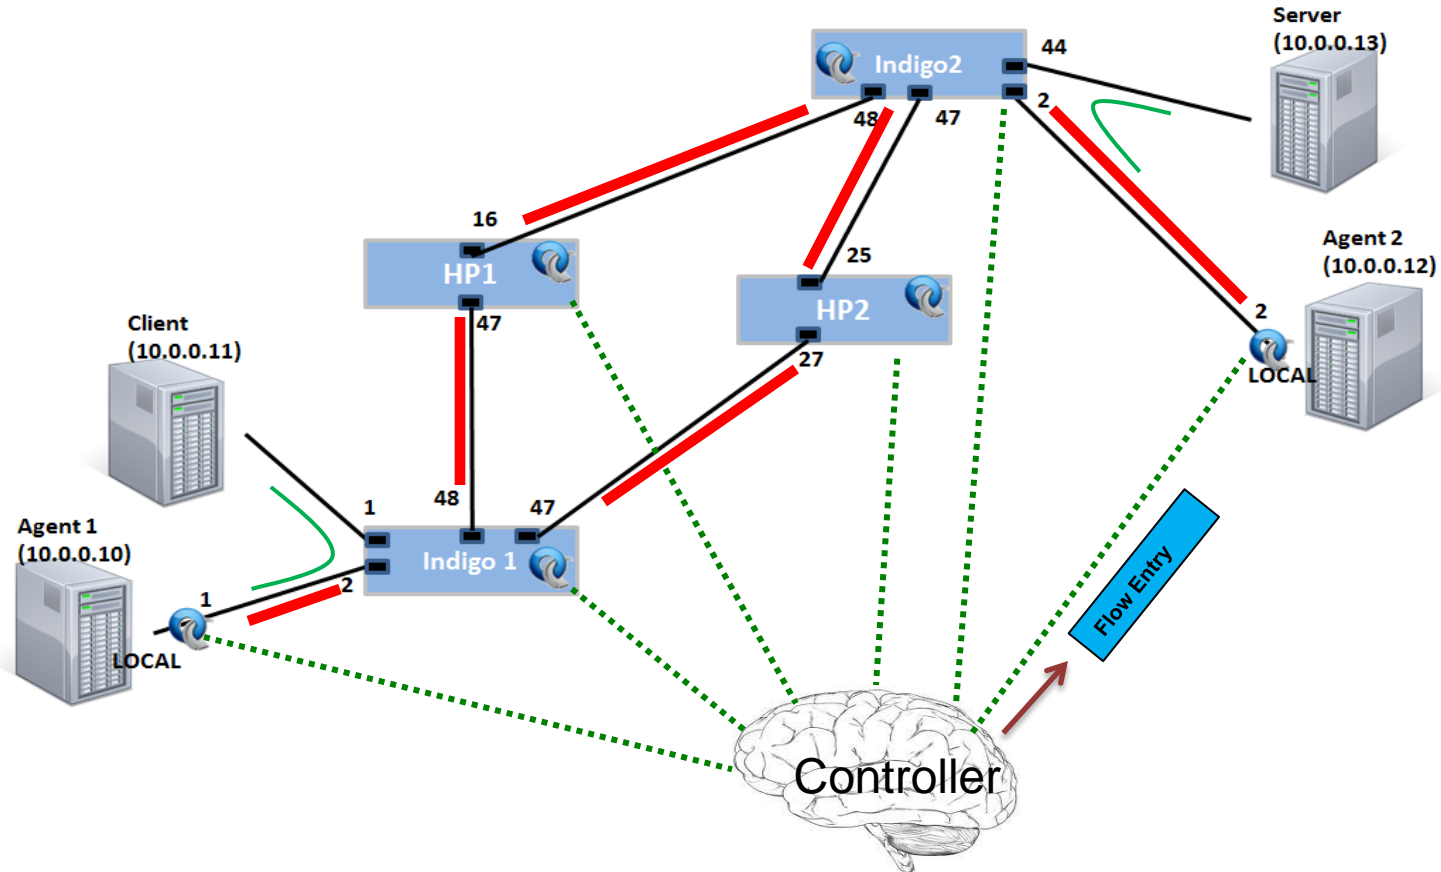
# Controller Handles packet

- ..... OpenFlow Control Plane
- Physical Connection



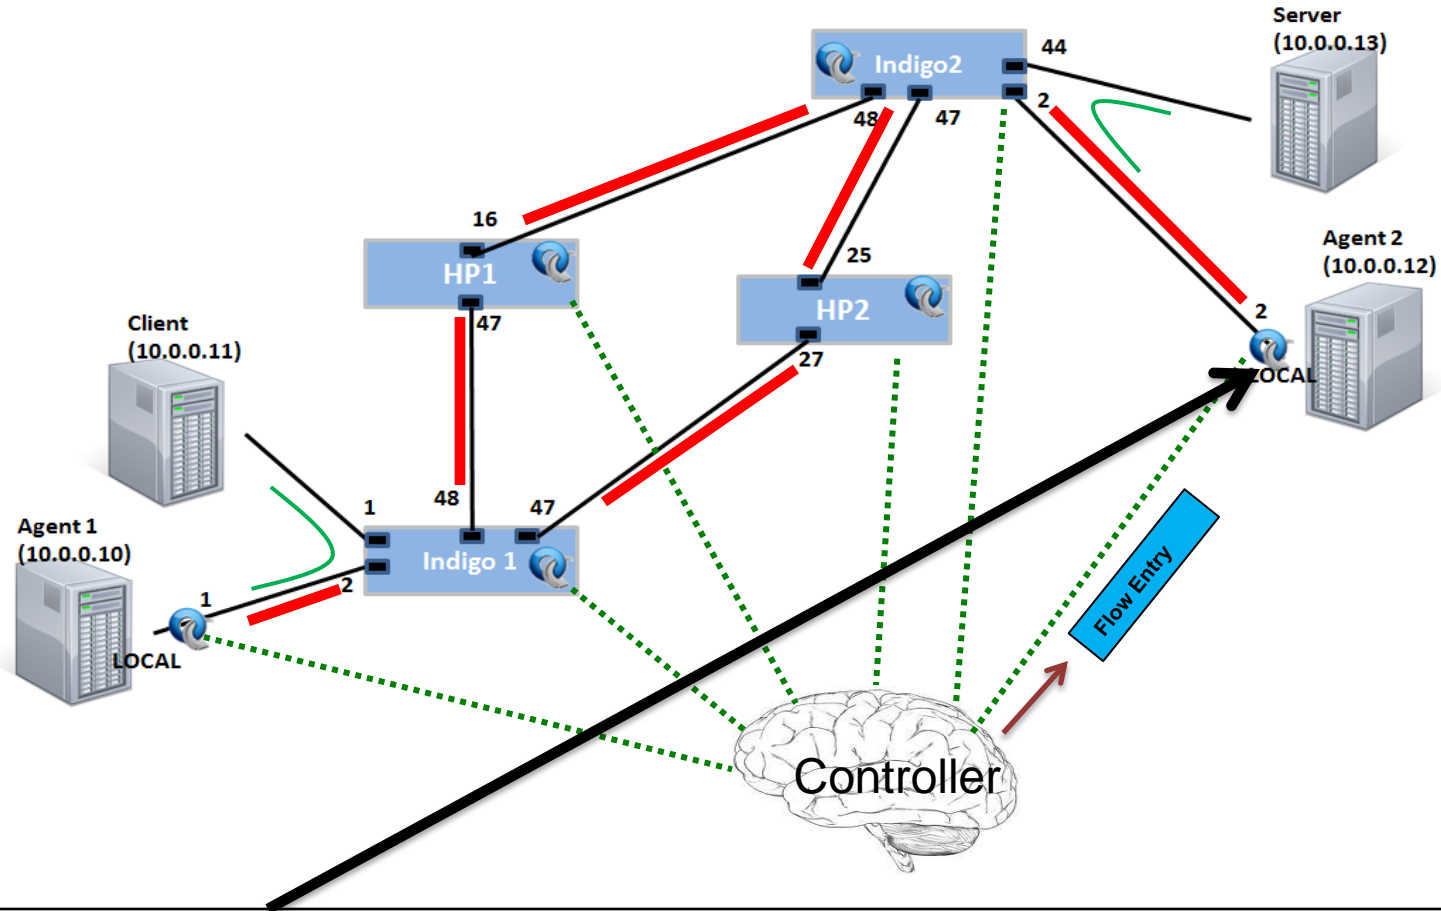
# Installs Flow Entry

- ..... OpenFlow Control Plane
- Physical Connection



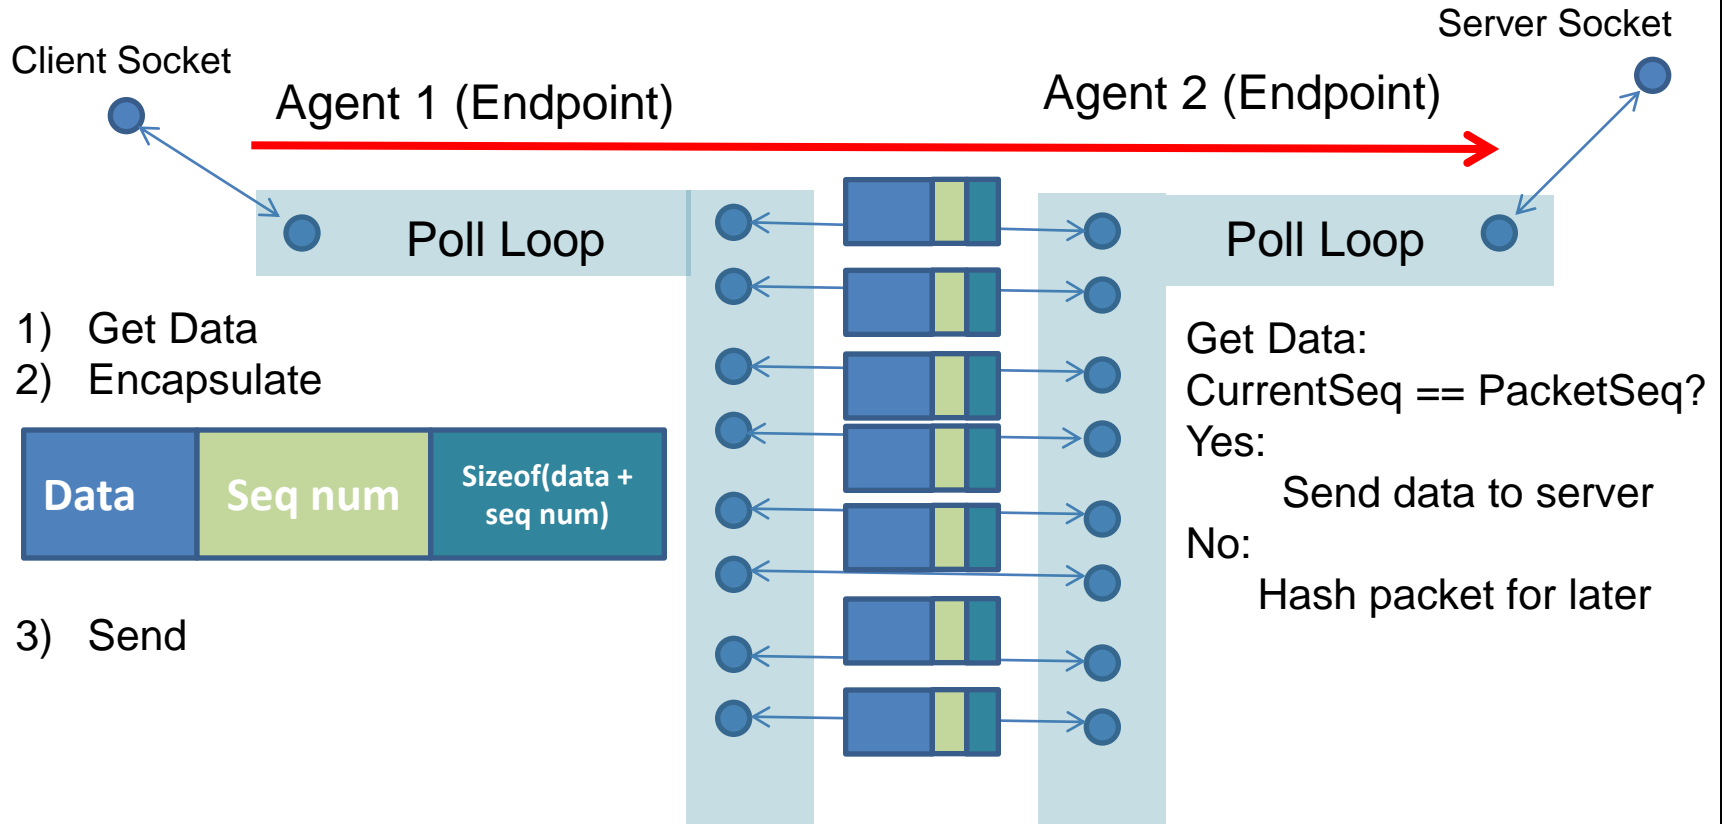
# Last Flow Entry for Agent 2 to Server

- ..... OpenFlow Control Plane
- Physical Connection



```
in_port=65534,vlan_tci=0x0000,dl_src=00:1b:21:6b:50:df,dl_dst=00:1b:21:6a:85:88,nw_src=10.0.0.12,nw_dst=10.0.0.13,nw_tos=0,
tp_src=47489,tp_dst=5003 actions=mod_dl_src:00:1f:29:32:91:99,mod_nw_src:10.0.0.11,output:2
```

# SOS Sockets



# SOS Transmission Protocol

- SOS uses Parallel TCP in order to relay data between agents.
- Exploits buffer limits by using multiple sockets.
- Sending window grows much faster using multiple sockets.
- Improves throughput in scenarios:
  - Lossy network
  - High Bandwidth Delay Product
  - Multiple paths present

# GENI

- Large multipath test bed across the US using Internet2 and National LambdaRail.
- Equipped with OpenFlow at the edge and in the core of the network.
- Allows full network visibility and control in core network.
- Series of different compute resources (dedicated, shared).
- Allows experimenters to test ideas on a real network
- Uses network slicing to allow multiple users





# GENI Results




Configuration	Latency (ms)	TCP (Mbps)	UDP (Mbps)	SOS (Mbps)
Path 1 (short)	54	295	952	869
Path 2 (long)	160	95	952	788
Multipath				703

\*Max average SOS throughput achieved when running for 3 minutes with varied number of sockets. (Could be better tuned to improve results)

# Test network


- In order to accurately measure network traffic a separate network was created and tc was used to emulate different network characteristics (latency, bandwidth, loss).



# TCP Buffer Space


- Main limit to TCP throughput in large bandwidth delay product network is lack of buffer space. Test below shows that sending window grows exponentially until buffer limit is reached.

**Sending Window Size vs Time with 600ms RTT**




# TCP Latency Throughput Growth Rate

- Latency affects the rate of which a TCP stream will grow since TCP window growth is updated based on RTT



# Parallel TCP Solution

- Using multiple TCP sockets allows the buffer size limit to be exploited
- TCP slow-start is faster
- Graph shows linear growth until congestion occurs



## Receiver Queue Benefits

- TCP suffers from head of line blocking
- Due to out of order arrivals at network layer
  - Retransmitted packets
  - Packet-level multipath routing
  - Route fluttering
  - And many other reasons
- Parallel TCP suffers from same issues in addition to at the application level
  - Out of order packet arrivals among TCP streams
- Receiver queue allows agent to receive data from streams while others are blocked
- Helps increase performance when:
  - Multiple paths of varied latency and bandwidth are used
  - Lossy network paths are present

# Latency Effect on Throughput with Receiver Queue

- Network setup with two 300Mbit paths
- Latency was then varied on one path and kept constant on the other
- Queue size represents the number of application level packets that can be stored

Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

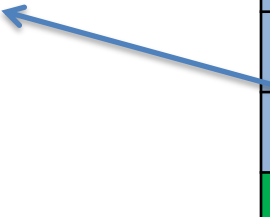
\*Note: when 1 stream is used it takes path of higher latency

Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# Receiver Queue Provides Little Benefit using Paths of Same Latency

- Receiver queue helps increase performance when multiple paths of varied latency are used
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking

0%  
Performance  
Improvement



Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

\*Note: when 1 stream is used it takes path of higher latency

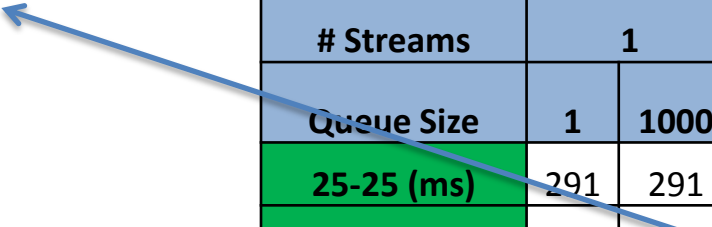
Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805



# Queue Yields Performance Increase When Path Latency Varies

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

12.5%  
Performance  
Improvement



Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

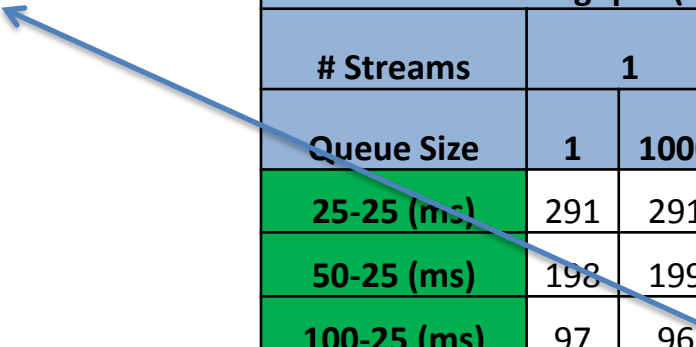
\*Note: when 1 stream is used it takes path of higher latency

Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# Queue Yields Performance Increase When Path Latency Varies

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

25%  
Performance  
Improvement



Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

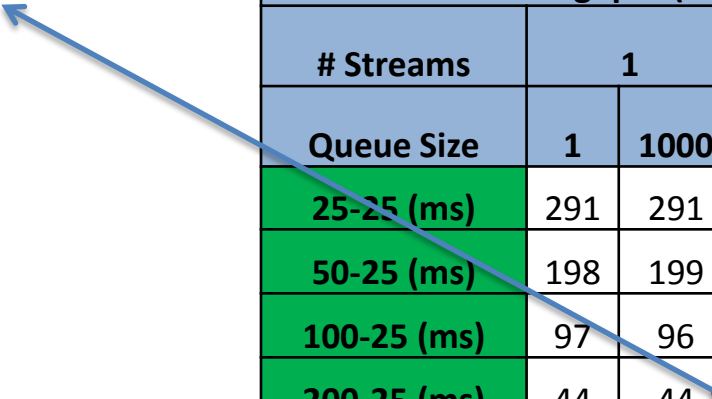
\*Note: when 1 stream is used it takes path of higher latency

Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# Queue Yields Performance Increase When Path Latency Varies

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

59.2%  
Performance  
Improvement



Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

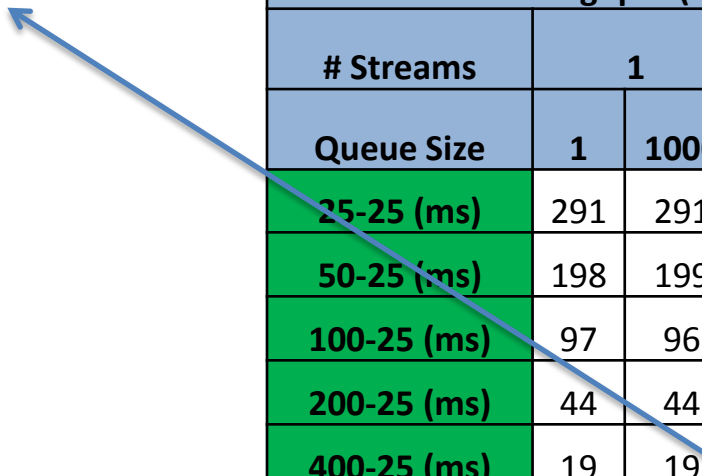
\*Note: when 1 stream is used it takes path of higher latency

Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# Queue Yields Performance Increase When Path Latency Varies

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

103.3%  
Performance  
Improvement



Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

\*Note: when 1 stream is used it takes path of higher latency

Average Queue Length (Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805


# Low Queue Utilization with Similar Path Latency

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

\*Note: when 1 stream is used it takes path of higher latency

Not very utilized  
with similar latency



Average Queue Utilization(Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# High Queue Utilization with Varied Path Latency

- Receiver queue helps increase performance when multiple paths of varied latency are used.
- Need to have enough queue space to buffer the difference in bandwidth delay product of paths in order to avoid blocking.

Throughput (Mbps) vs Number of streams								
# Streams	1		2		4		6	
Queue Size	1	1000	1	1000	1	1000	1	1000
25-25 (ms)	291	291	580	580	580	584	580	585
50-25 (ms)	198	199	432	486	567	562	569	572
100-25 (ms)	97	96	304	380	452	478	524	541
200-25 (ms)	44	44	157	250	309	336	371	390
400-25 (ms)	19	19	60	122	139	209	183	278

More utilized when varied path latency

\*Note: when 1 stream is used it takes path of higher latency

Average Queue Utilization(Bytes) Comparing Paths of Varied Latency (2 Streams)									
25-25(ms)		50-25(ms)		100-25(ms)		200-25(ms)		400-25(ms)	
.00176	.00142	.00101	.04039	.00115	.32321	.01184	.623	.1411	.4805

# Queue Size Effect with Multiple Paths of Varied Bandwidth

Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.

# Receiver Queue Provides Little Benefit using Paths of Same Bandwidth

.51%  
Performance  
Improvement

Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

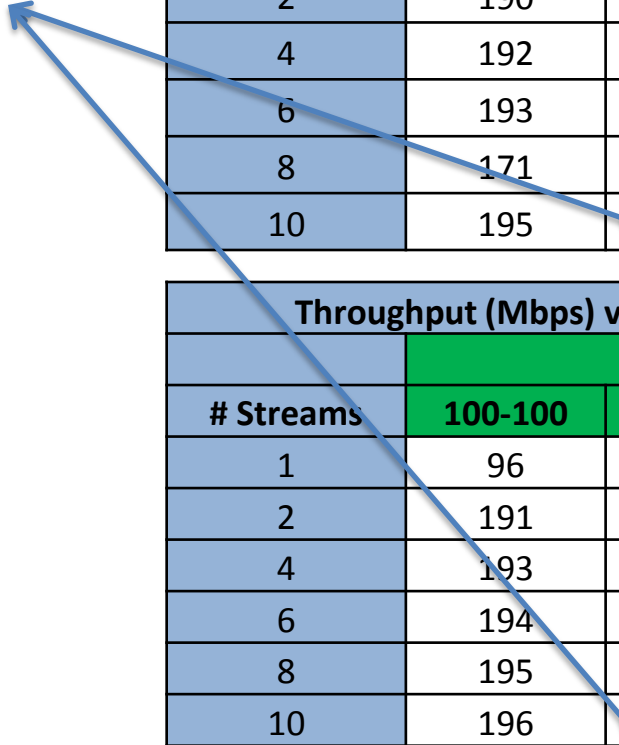
Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.



# Receiver Queue Provides Large Benefit using Paths of Same Bandwidth

47%  
Performance  
Improvement



Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.

# Receiver Queue Provides Large Benefit using Paths of Same Bandwidth

98.45%  
Performance  
Improvement

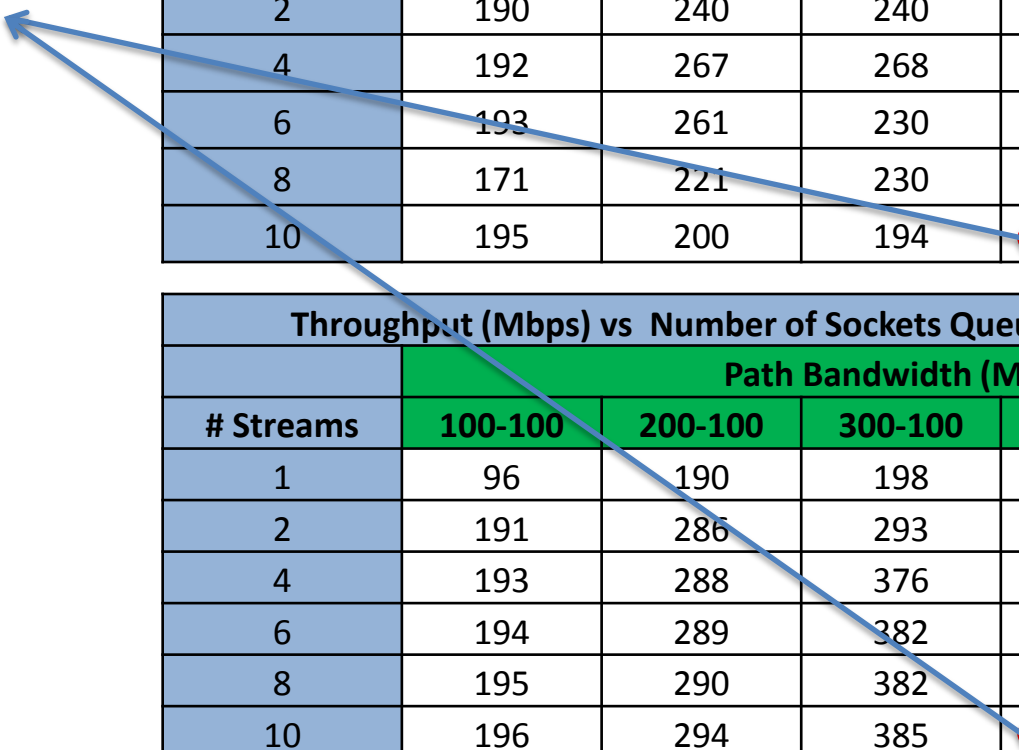
Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.

# Receiver Queue Provides Large Benefit using Paths of Same Bandwidth

98.73%  
Performance  
Improvement



Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.

# Receiver Queue Provides Large Benefit using Paths of Same Bandwidth

155.4%  
Performance  
Improvement

Throughput (Mbps) vs Number of Sockets Queue length 1					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	95	190	199	198	199
2	190	240	240	238	238
4	192	267	268	272	267
6	193	261	230	198	279
8	171	221	230	260	239
10	195	200	194	236	204

Throughput (Mbps) vs Number of Sockets Queue length 10,000					
	Path Bandwidth (Mbit)				
# Streams	100-100	200-100	300-100	400-100	500-100
1	96	190	198	199	199
2	191	286	293	293	293
4	193	288	376	413	425
6	194	289	382	460	499
8	195	290	382	472	547
10	196	294	385	469	521

\* When 1 stream is used, it runs over the path of higher available bandwidth.

# Lossy Link

- SOS is able to overcome a lossy link using multiple streams coupled with user space Queuing.

1% Lossy Path vs Nonlossy Path Throughput (Mbps)				
# Sockets	Lossy Path Throughput (Queue Size 1)	Lossy Path Throughput (Queue Size 1,000)	Nonlossy Path Throughput Mbps (Queue Size 1)	Nonlossy Path Throughput Mbps (Queue Size 1,000)
1	16	29	805	897
2	30	32	902	877
4	51	90	900	904
6	93	110	884	899
8	93	124	902	902
10	121	136	883	888
20	135	149	902	904
30	213	316	905	906
40	251	484	902	904
50	334	632	903	902
60	348	740	905	901
70	392	796	906	897
80	415	865	905	897
90	404	867	902	894
100	448	874	902	892

# Lossy Link Obtains High Throughput using Many Streams

- SOS is able to overcome a lossy link using multiple streams coupled with user space Queuing.

2,913.8%  
Performance  
Improvement

1% Lossy Path vs Nonlossy Path Throughput (Mbps)				
# Sockets	Lossy Path Throughput (Queue Size 1)	Lossy Path Throughput (Queue Size 1,000)	Nonlossy Path Throughput Mbps (Queue Size 1)	Nonlossy Path Throughput Mbps (Queue Size 1,000)
1	16	29	805	897
2	30	32	902	877
4	51	90	900	904
6	93	110	884	899
8	93	124	902	902
10	121	136	883	888
20	135	149	902	904
30	213	316	905	906
40	251	484	902	904
50	334	632	903	902
60	348	740	905	901
70	392	796	906	897
80	415	865	905	897
90	404	867	902	894
100	448	874	902	892

# Lossy Link Achieves Similar throughput as Nonlossy Link

- SOS is able to overcome a lossy link using multiple streams coupled with user space Queuing.

2.6%  
Performance  
Difference

1% Lossy Path vs Nonlossy Path Throughput (Mbps)				
# Sockets	Lossy Path Throughput (Queue Size 1)	Lossy Path Throughput (Queue Size 1,000)	Nonlossy Path Throughput Mbps (Queue Size 1)	Nonlossy Path Throughput Mbps (Queue Size 1,000)
1	16	29	805	897
2	30	32	902	877
4	51	90	900	904
6	93	110	884	899
8	93	124	902	902
10	121	136	883	888
20	135	149	902	904
30	213	316	905	906
40	251	484	902	904
50	334	632	903	902
60	348	740	905	901
70	392	796	906	897
80	415	865	905	897
90	404	867	902	894
100	448	874	902	892

# Sending/Receiving Buffer Size Effect

- The number of bytes agent should try to send/receive each time was found to greatly affect performance.
  - `Send(fd, buf, sizeof(buf), 0);`
- Initial thought was due to overhead added by additional headers.


27.7%  
Performance  
Improvement

Varying Buffer Size Results (100ms RTT)					
Buffer Size (Bytes)	512	1024	2048	4096	8192
Throughput (Mbps)	155	156	169	188	198
# Application Level Packets	2311930	1152818	622046	345238	182105
% Overhead	1.56	0.78	0.39	0.195	0.097
Average Packet Size Send/Received (Bytes)	511.99	1023.99	2047.99	4095.97	8195.92



# SOS Overhead

- Addition 8 bytes was added in order to breakup and reassemble data
- Size of packet and sequence number
- $8/512 = 1.56\%$
- $8/8196 = .0976\%$
- Not always true do to amount available to read



## Accounting for overhead:

Sends Block size – Header packets  
 $(2311930 * (511.99 - 8) * 8) / (60 * 10^6)$   
: **155.35Mbps**

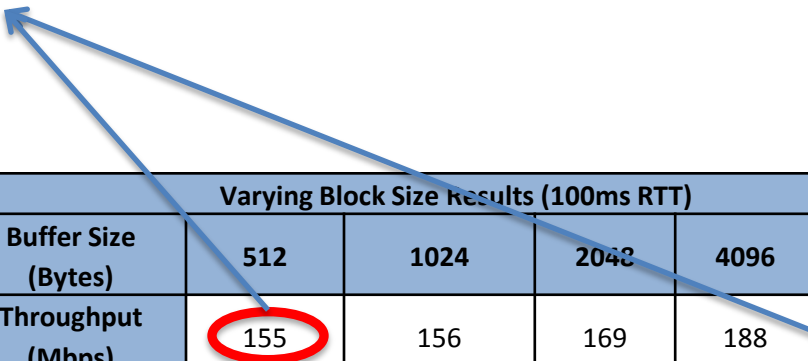
Assuming no overhead  
 $(2311930 * (511.99) * 8) / (60 * 10^6)$   
: **157.82Mbps**

: **2.46Mbps Lost due to overhead**

Varying Block Size Results (100ms RTT)					
Block Size (Bytes)	512	1024	2048	4096	8192
Throughput (Mbps)	155	156	169	188	198
# Application Level Packets	2311930	1152818	622046	345238	182105
% Overhead	1.56	0.78	0.39	0.195	0.098
Average Packet Size Send/Received (Bytes)	511.99	1023.99	2047.99	4095.97	8195.92

# Missing Throughput

- 2.46Mbps Lost due to overhead
- Though,  $198 - 157.82 = 40.18\text{Mbps}$  Where did this go?



Buffer Size (Bytes)	512	1024	2048	4096	8192
Throughput (Mbps)	155	156	169	188	198
# Application Level Packets	2311930	1152818	622046	345238	182105
% Overhead	1.56	0.78	0.39	0.195	0.098
Average Packet Size Send/Received (Bytes)	511.99	1023.99	2047.99	4095.97	8195.92

## Maybe due to TCP Overhead?

- Lower Average Segment Size when using smaller buffer
- Requires more packets (+54 Byte header per packet)
- Values obtained via tcptrace which analysis's the pcap trace

Average Segment Size vs Block Size (100ms RTT)					
Buffer Size (bytes)	512	1024	2048	4096	8192
Average Segment Size (bytes)	1493	1494	2107	4193	8354
# Packets Sent	792,416	789,710	604,612	337,199	178,640



**343.58% More packets sent, though 27.7% slower!**

## Maybe due to TCP Overhead?

- Lower Average Segment Size when using smaller buffer
- Requires more packets (+54 Byte header per packet)
- Values obtained via tcptrace which analysis's the pcap trace

Average Segment Size vs Block Size (100ms RTT)					
Buffer Size (bytes)	512	1024	2048	4096	8192
Average Segment Size (bytes)	1493	1494	2107	4193	8354
# Packets Sent	792,416	789,710	604,612	327,199	178,640



Why is this bigger than my 1500MTU?

# TCP Segmentation Offload (TSO)

- Offload Segmentation to NIC
  - Can pass down 64KB of data to NIC which breaks into MTU chunks
  - Drastically reduces CPU time needed to transmit data
  - Segment Size does come into play with high throughput

Average Segment Size vs Block Size (100ms RTT)					
Block Size (bytes)	512	1024	2048	4096	8192
Throughput (Mbps)	155	157	157	158	159
Average Segment Size (bytes)	1447	1447	1447	1447	1447

Average Segment Size vs Block Size (0ms RTT)					
Block Size (bytes)	512	1024	2048	4096	8192
Throughput (Mbps)	464 (816)	587 (824)	826 (852)	863 (893)	870 (909)
Average Segment Size (bytes)	1335	1392	1447	1447	1447

# Improving Performance

- No one magic number for amount for number of sockets to use.
- What information can SOS use to improve performance?
- Queue Utilization, Average sent/received block size
- Rexmt is the number of retransmitted data

# Streams	Throughput (Mbps)	Average Received Block Size (Bytes)	Average Queue Utilization	Rexmt (MB)	Sent Data in 5 minutes (GB)
5	597	8180.2	.07026	0	20.78
100	512	7217.98	.62778	3.4	17.94

# Improving Performance

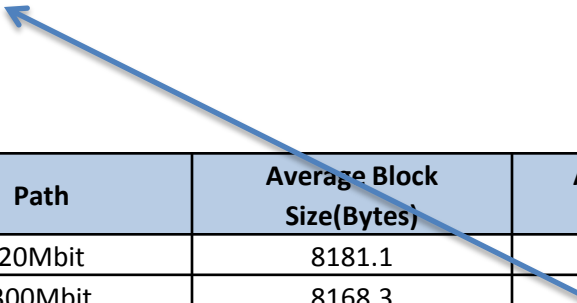
- Using the average receive block size and average queue utilization agent could detect over utilization and disable sockets to improve performance

# Streams	Throughput (Mbps)	Average Received Block Size (Bytes)	Average Queue Utilization	Rexmt (MB)	Sent Data in 5 minutes (GB)
5	597	8180.2	.07026	0	20.78
100	512	7217.98	.62778	3.4	17.94

**Signs of over utilization**

# Improving Performance

- Multipath example with varied available bandwidth.
- 20 streams used which achieved throughput 111Mbps
- Agent spent most of time blocking on streams that traversed 300Mbit path.



Path	Average Block Size(Bytes)	Average Queue Utilization
20Mbit	8181.1	.13529
300Mbit	8168.3	.92177

- Agent could detect this and dynamically alter the number of streams that traversed each path to improve performance.
- Not Implemented nor evaluated.



# Conclusion

---

- SOS shows how network providers can decouple users choice of protocols to network providers choice of protocols.
- Alleviates complexities for end user since no modifications are required.
- Throughput improvements are shown over high Delay Bandwidth Product Networks and Lossy links
- Receiver Queue significantly helps to improve throughput results in multipath scenarios.

# Future Work

---

- Improve SOS to dynamically alter number of sockets used in order to optimize throughput
- Develop other services such as client mobility, delay tolerant network services, and increased security could be provided via similar methods